

# Microservices

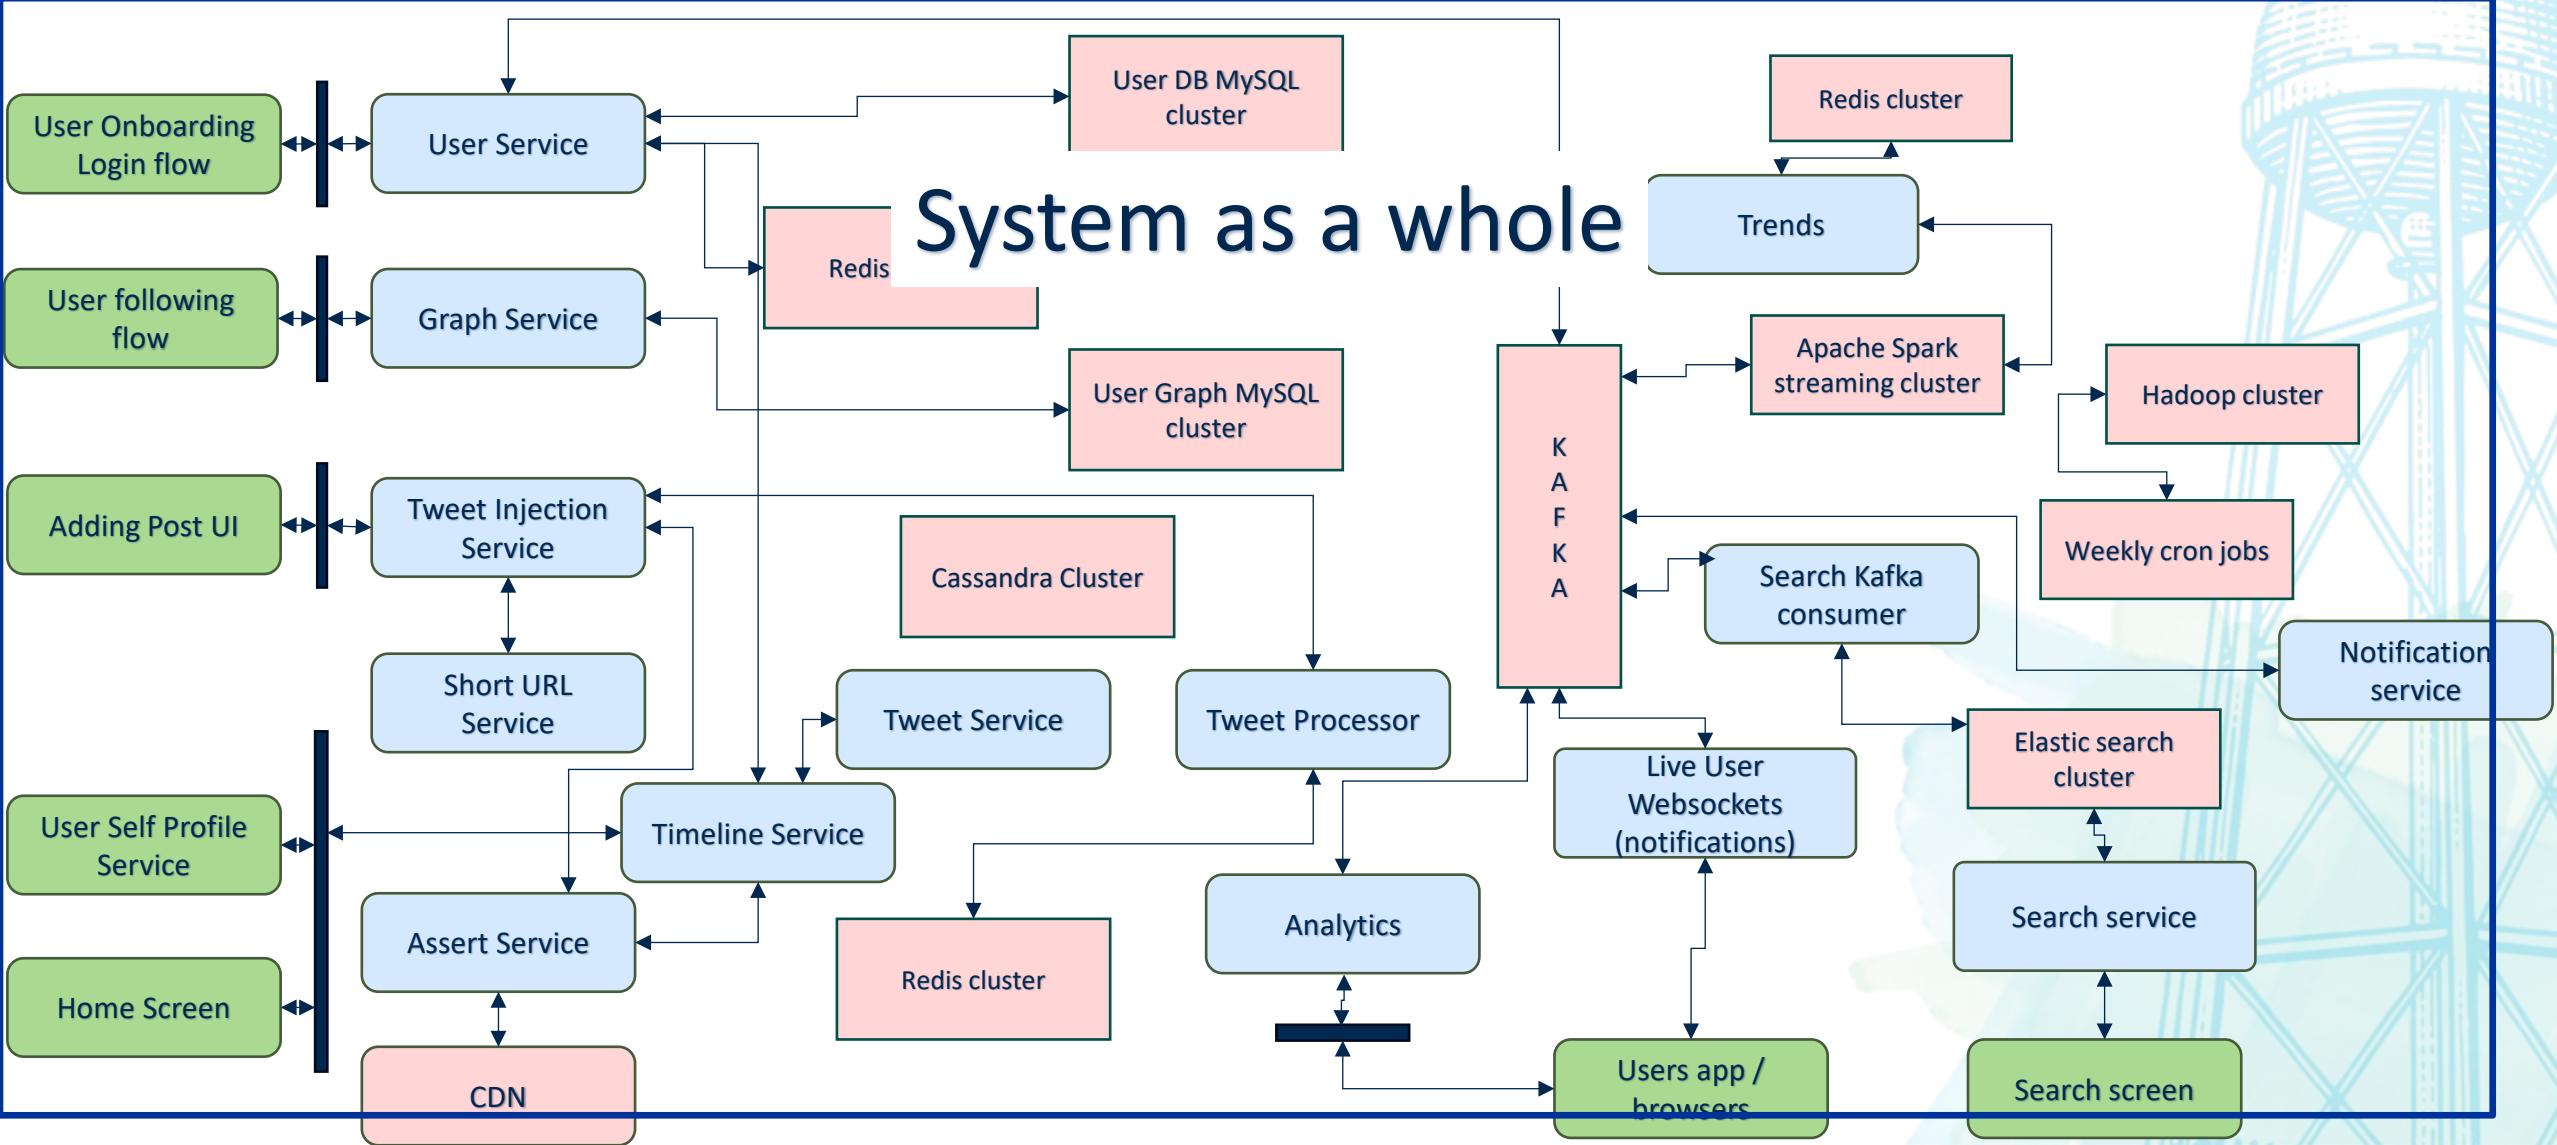
Tapti Palit

# Outline

- Service boundaries
- Communication (sync vs. async)
- Data per service
  - Logs (events)
  - LSM (state)
  - In-memory cache

# Software architecture

## System as a whole

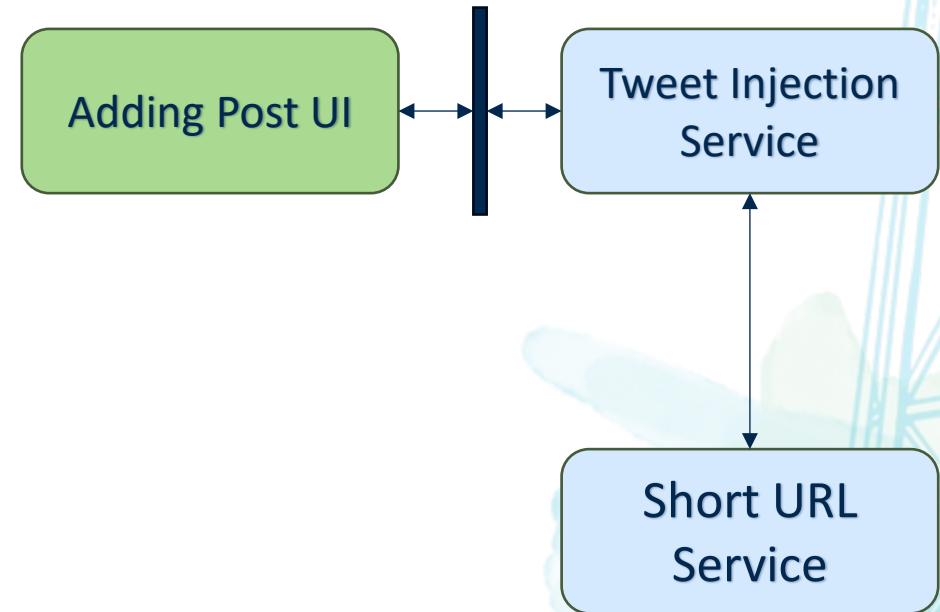


# Some software architecture concerns

- Encapsulation
- Scalability
- Resilience
- Performance
- Ease of deployability
- Security (ECS 153)

# Encapsulation

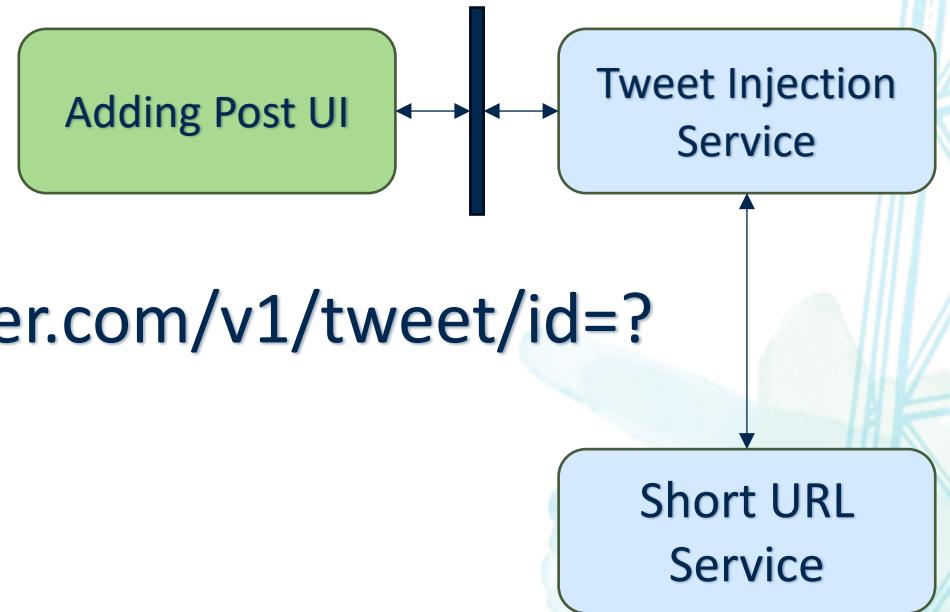
- At the component level
  - Does TweetInjectionService need to know anything about the *internals* of ShortURLService?



# Encapsulation

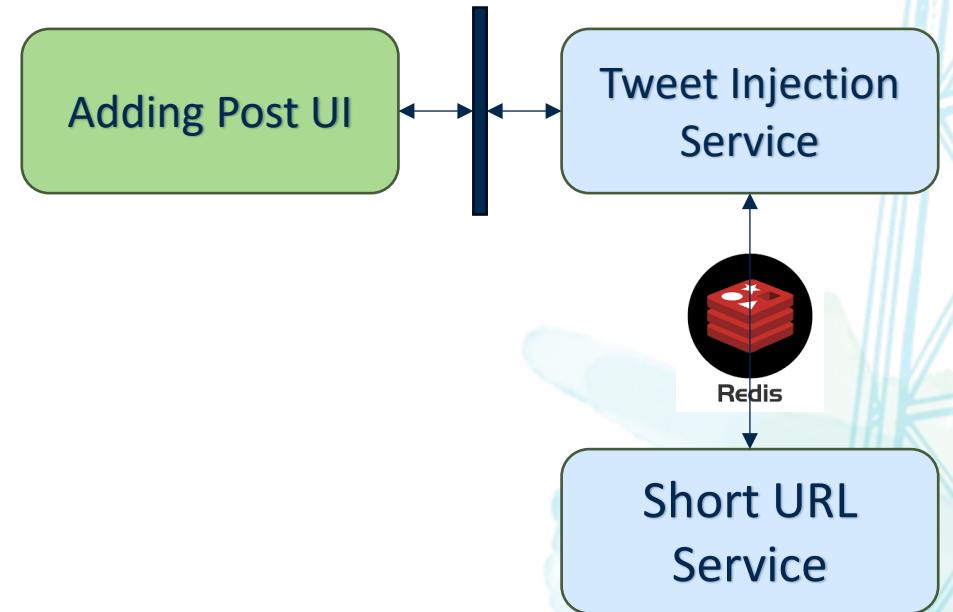
- What interface does ShortURLService provide to TweetInjectionService?

<https://twitter.com/v1/tweet/id=?>



# Encapsulation

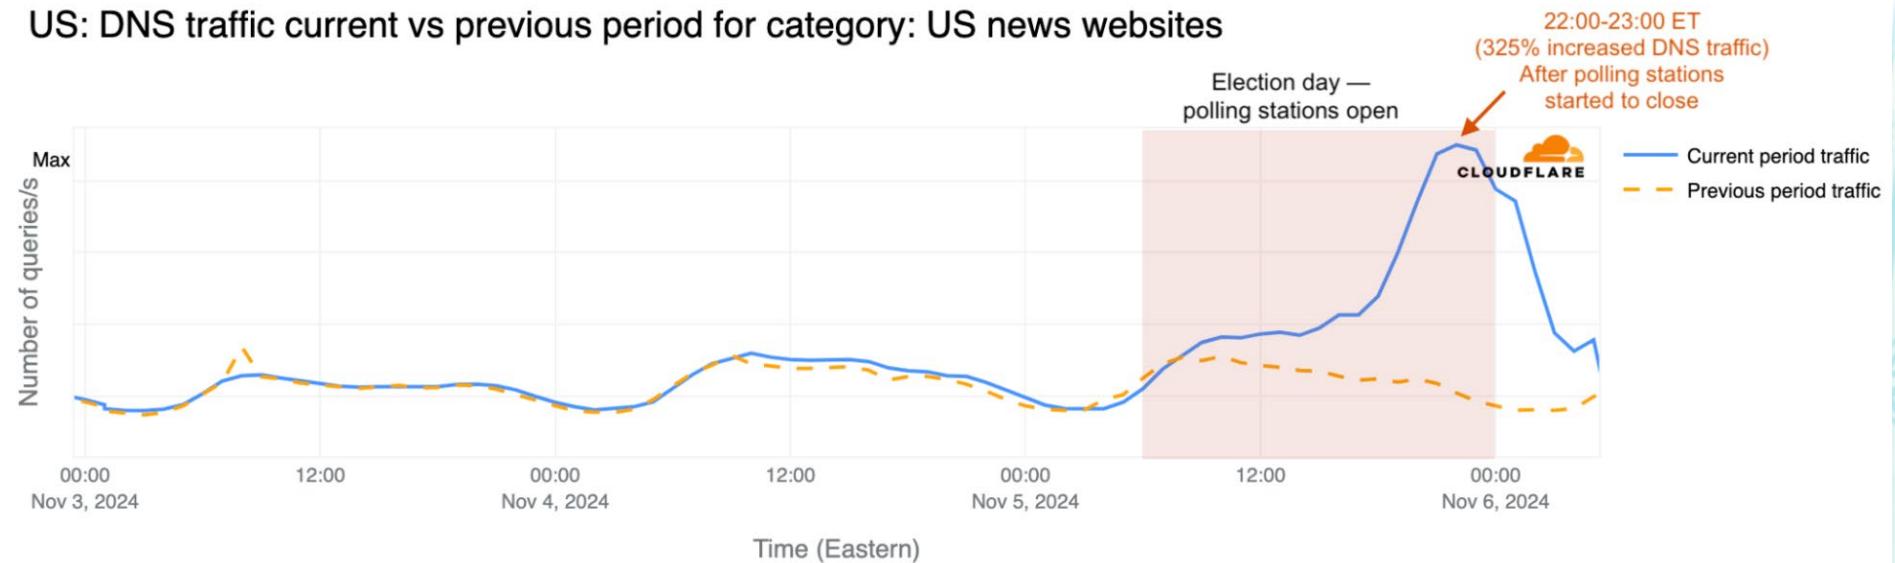
- Typically run on different machines
  - How do they **communicate?**
  - Directly call each other or is there a mediator?
    - Different tradeoffs



# Scalability

- Website traffic is not constant
- Can spike due to planned events
  - Product launch
  - US elections

<https://blog.cloudflare.com/exploring-internet-traffic-shifts-and-cyber-attacks-during-the-2024-us-election/>



# Scalability

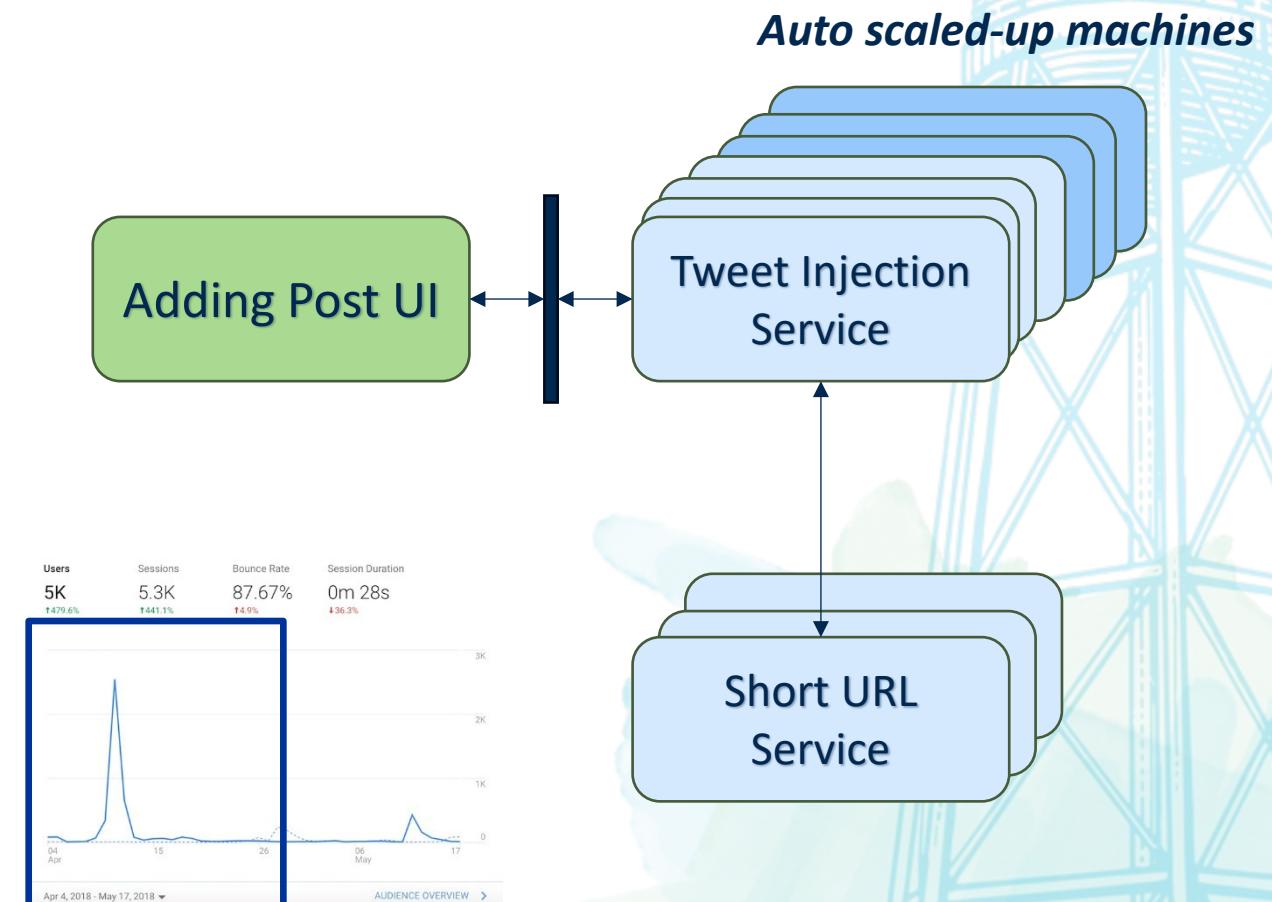
- Unplanned events
  - Post goes viral
  - Architecture should *scale* to handle such events

<https://www.residualthoughts.com/2018/05/20/traffic-data-from-a-viral-post/>



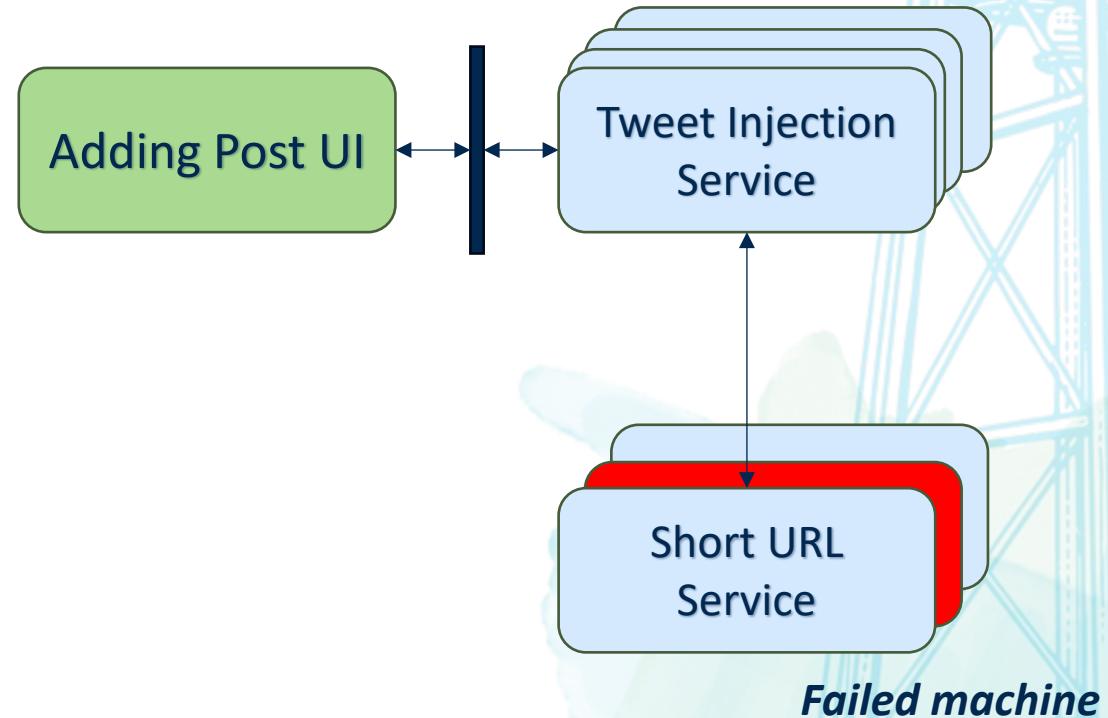
# Scalability

- Services run on more than one machine
- Autoscaling for when traffic reaches threshold
  - System automatically brings up new “machines”
  - Available in popular cloud deployments
    - Amazon AWS, Digital Ocean, Google Cloud



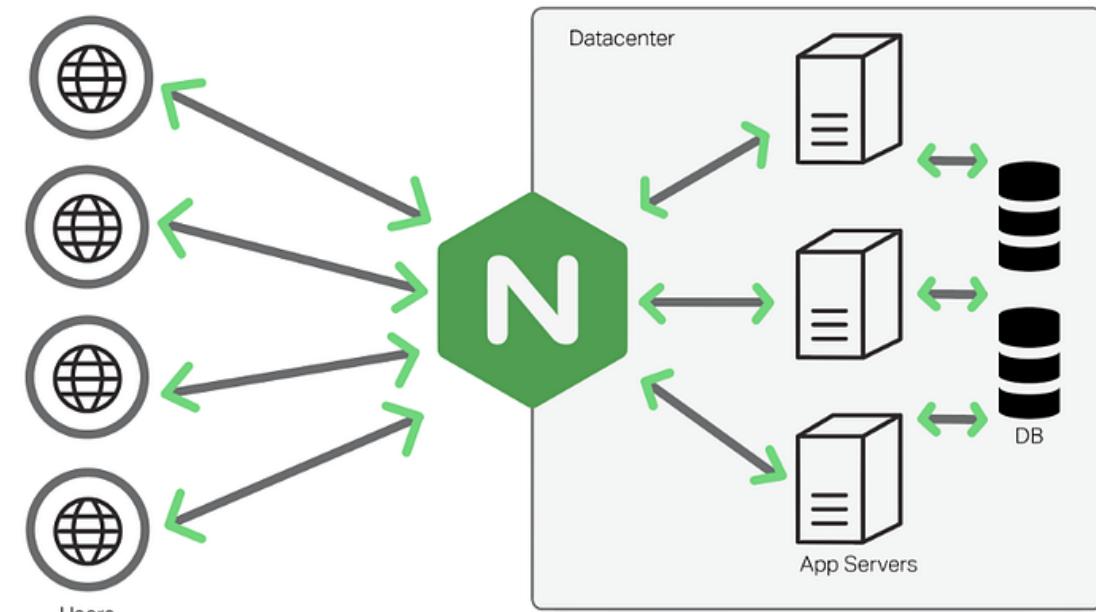
# Resilience

- Is the system resilient to failures?
  - If a single component fails does the entire system fail?



# Resilience

- Fault tolerance
  - Distribute traffic across multiple instances (load balancers)
  - Load balancer detects the failures and stops sending requests to faulting nodes



*Nginx load balancer*

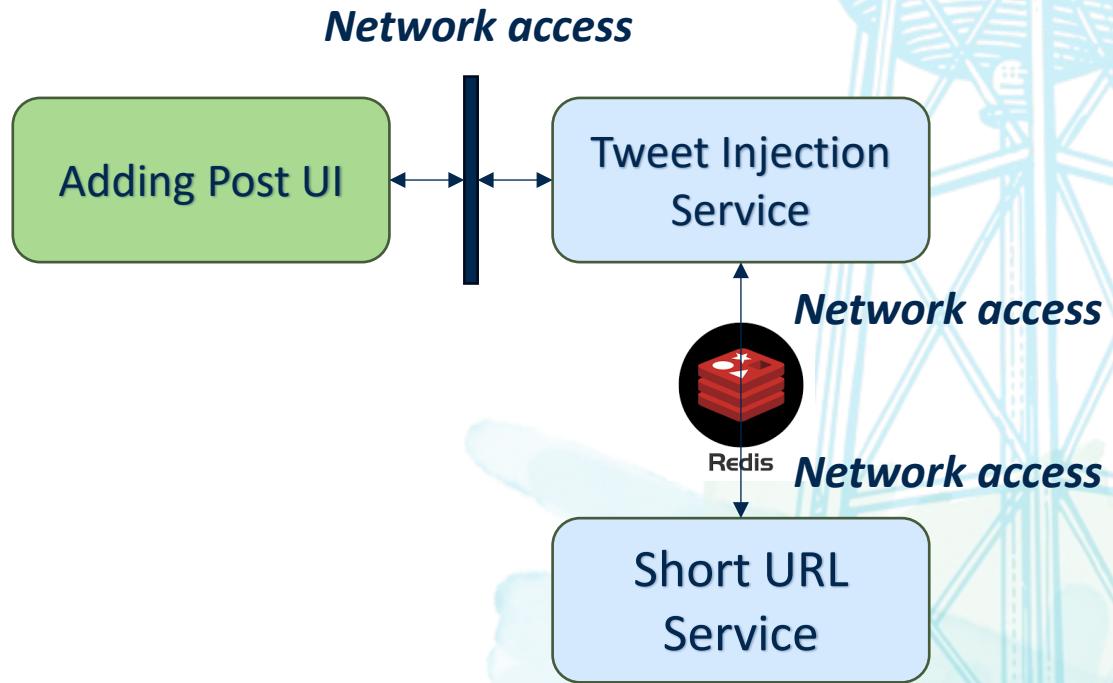


HAProxy



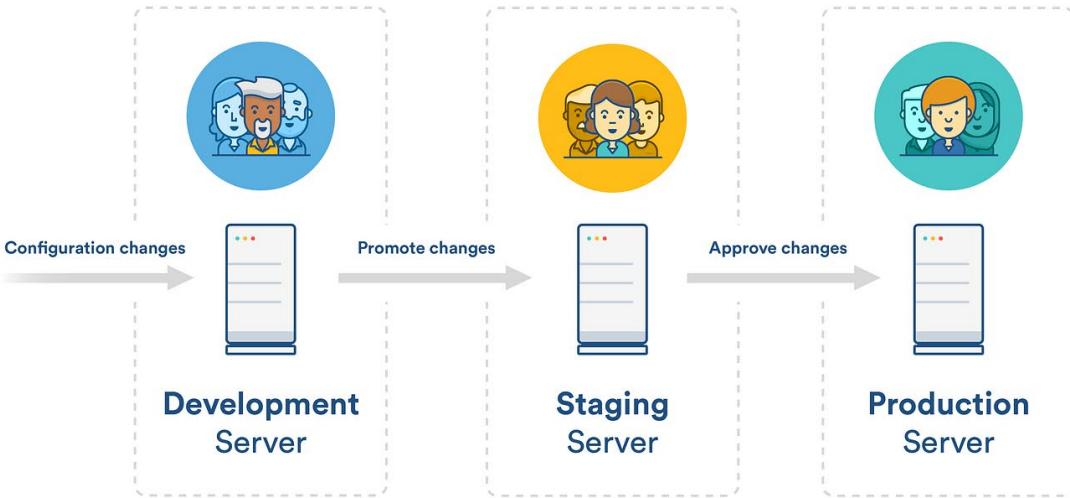
# Performance

- Does the architecture itself cause any performance bottlenecks?
- Modularity reduces performance
- How to limit performance overhead?



# Deployability

- Can individual components be individually deployed?
  - Do I need to deploy TweetService to deploy ShortURLService?
- Can deployment steps be automated?



# Dimensions of software architecture styles

- System organization
- Data organization
- Communication architecture

# Dimensions of software architecture styles

- System organization
  - Monolithic and microservices
- Data organization
- Communication architecture

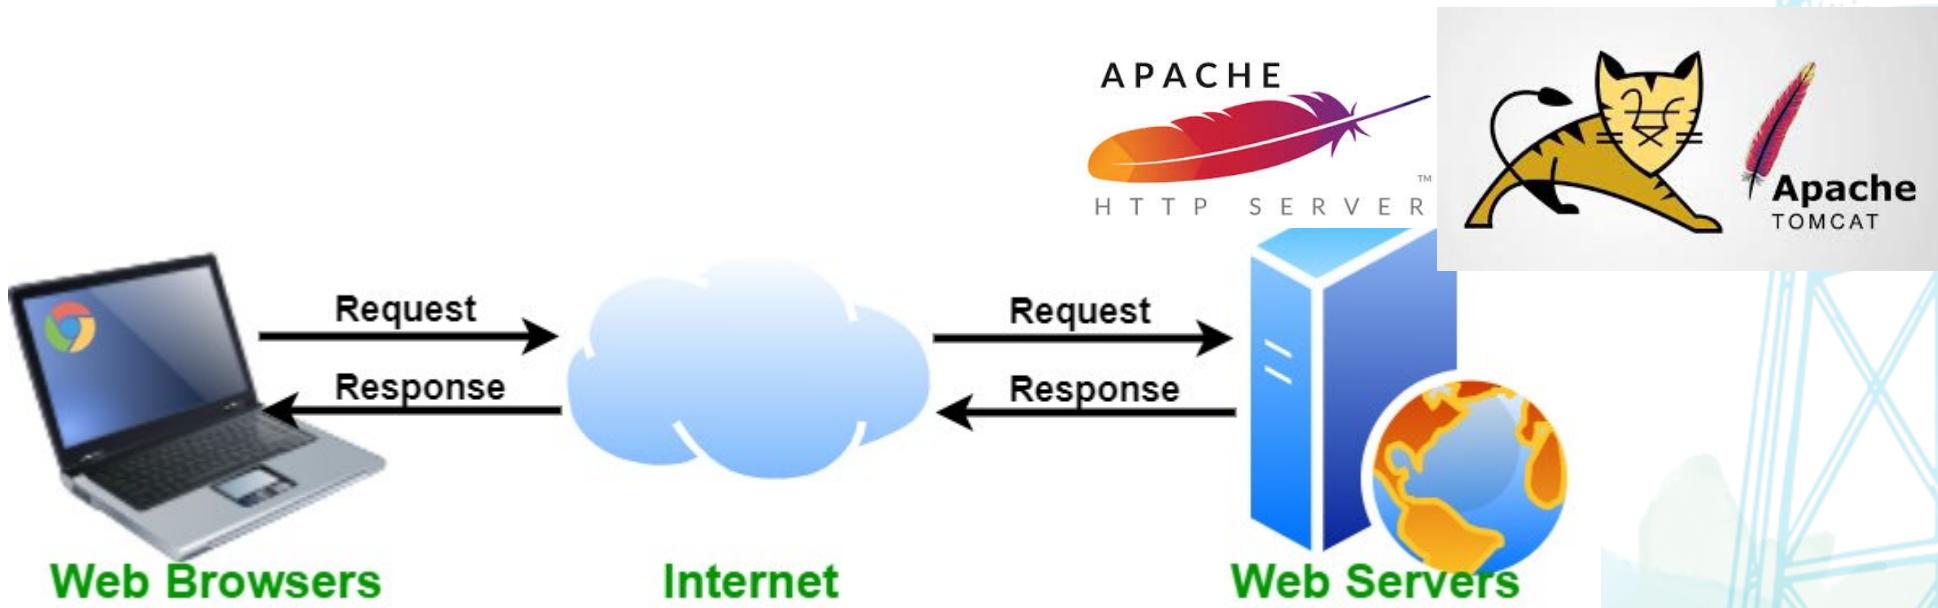
# Dimensions of software architecture styles

- System organization
  - Monolithic and microservices
- Data per service
  - LSM (state)
  - Logs (events)
  - In-memory caches
- Communication architecture

# Dimensions of software architecture styles

- System organization
  - Monolithic and microservices
- Data **per service**
  - LSM (state)
  - Logs (events)
  - In-memory caches
- Communication architecture **for microservices**
  - Synchronous (RPC)
  - Asynchronous (Message queuing, pub-sub)
  - Event-driven (Kafka – next module)

# Client server architecture



# Pros and cons

- Benefits
  - Centralized management of data: all data is on the server and can be easily secured
    - No complex data flows
  - Separation between client and server business logic
- Disadvantages
  - Single point of failure
  - Poor scaling

# Microservices

- An approach to developing a single application as a suite of small services, each running in its own process and communicating with lightweight mechanisms, often an HTTP resource API <https://martinfowler.com/articles/microservices.html>
- Independently deployable by automated processes
- Bare minimum centralized management
- Smart endpoints connected by “dumb” pipes

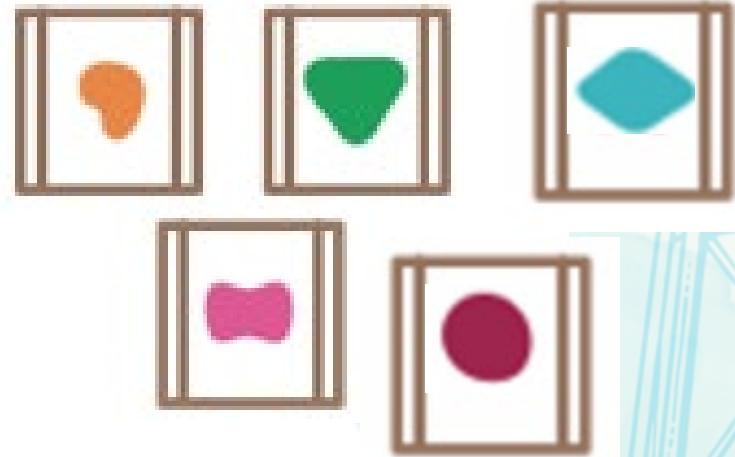
# Microservices overview

- An architectural style for building distributed systems
- Applications are divided into small, independent services



*Monolithic application*

*(Multiple components running in same process on same machine)*

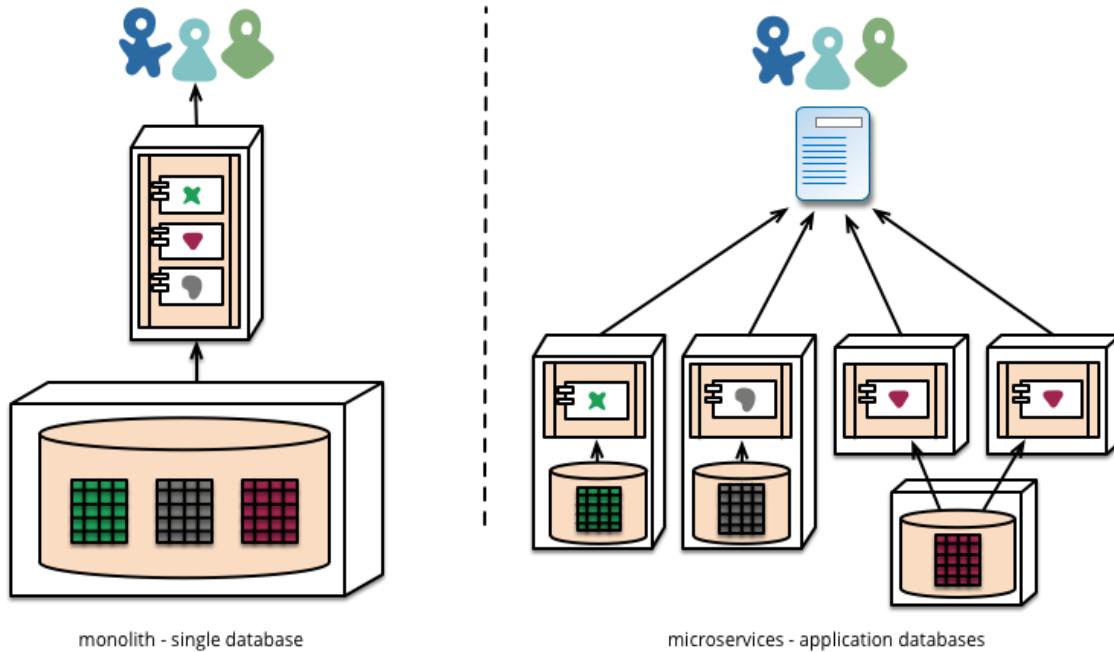


*Microservices*

*(Each component runs in its own process  
on potentially its own machine)*

# Monolithic to microservices

- Each microservice can have its own database – ***shared nothing architecture***



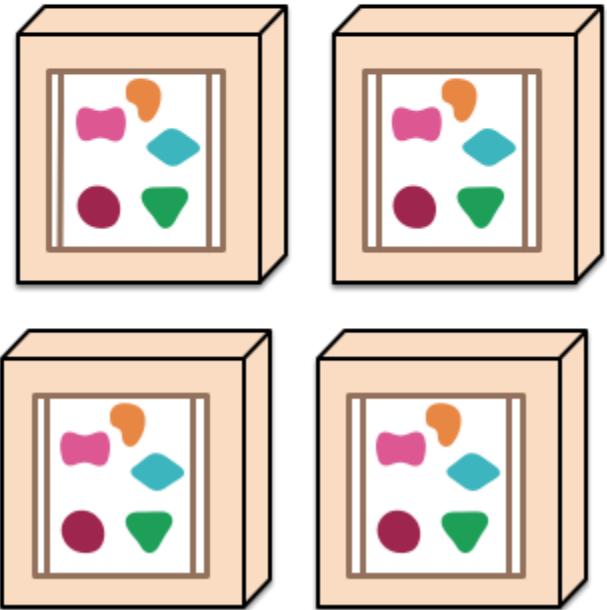
<https://martinfowler.com/articles/microservices.html>

# Microservices provide finer scaling control

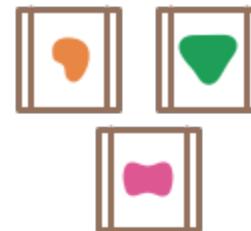
*A monolithic application puts all its functionality into a single process...*



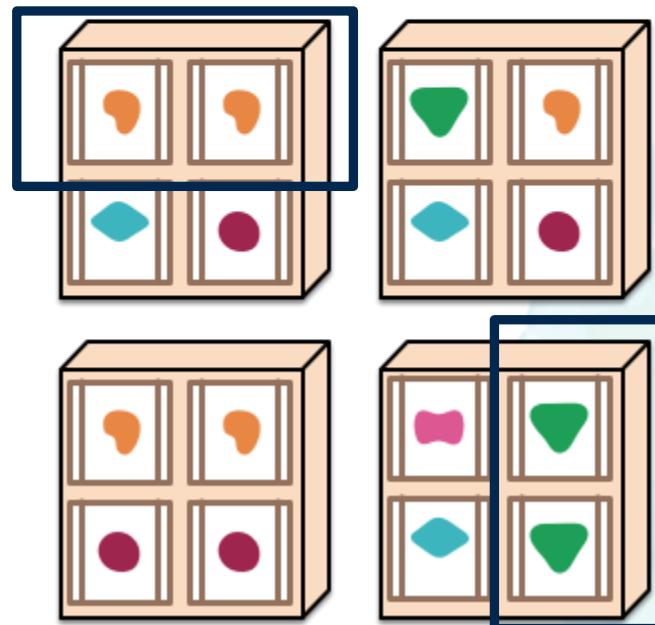
*... and scales by replicating the monolith on multiple servers*



*A microservices architecture puts each element of functionality into a separate service...*



*... and scales by distributing these services across servers, replicating as needed.*



***Finer control over scaling***

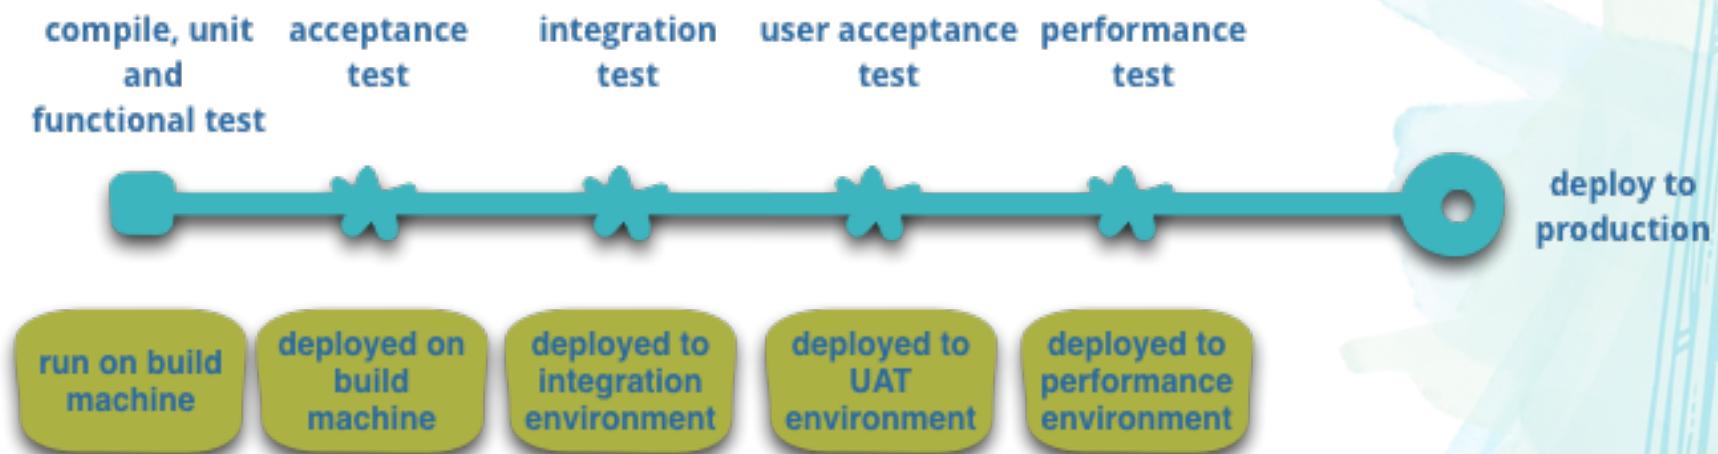
<https://martinfowler.com/articles/microservices.html>

# Microservices provide finer scaling control

- Monolithic application can only be deployed as a whole
  - Therefore, can only be scaled as a whole
- Microservices are individually deployable
  - Therefore, they provide finer scaling control

# Continuous deployment using microservices

- Microservices simplify automated deployment
- E.g., pushing to the main branch in GitHub triggers GitHub Actions that
  - Automatically builds
  - Execute unit-tests
  - Automatically deploy to integration environment (and then production environment)



# Pros and cons

- Pros
  - Strong encapsulation and modularity
  - Better reusability
  - Each microservice can be scaled independently (*More on this later*)
  - Each microservice can be written in its own programming language
  - Fault isolation
  - Supports CI/CD (easier to deploy microservices than monolithic services)
- Cons
  - Higher complexity
  - Debugging complex interactions is harder
  - Network overhead

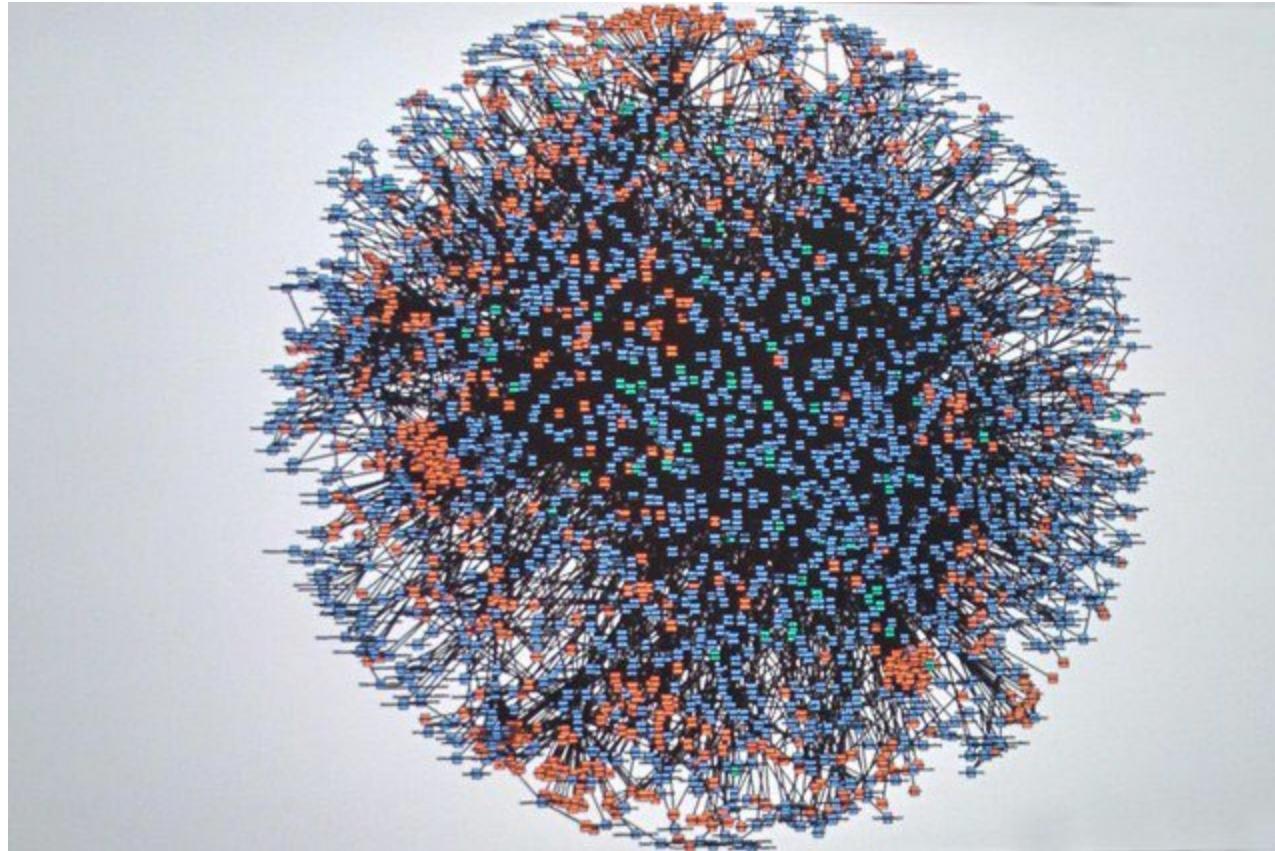
# Microservices at Uber (2019)



<https://x.com/msuriar/status/1110244877424578560>

# Microservices at Amazon (2008)

- Code-named “Deathstar”



<https://x.com/Werner/status/741673514567143424>

# Spring Boot Overview

- Framework for creating RESTful microservices
- Reduces boilerplate configuration code
- Embedded server (Tomcat/Jetty)
- Simplifies microservice creation through annotations
- Built-in support for REST APIs

# RESTful microservices with Spring Boot

- Create classes that can act as REST endpoints
- Uses annotations to denote REST endpoint URLs
  - Allows complete decoupling from the boilerplate code
- Types of requests
  - @GetMapping, @PostMapping, @PutMapping, and so on... for all HTTP methods
- @PathVariable – extract variable from GET request
- @RequestBody – extract the post request body

```
class MyRequest {  
    private String postDate;  
    private String postContent;  
    // ... Getters and setters  
}  
  
@RestController  
@RequestMapping("/myservice")  
public class MyController {  
    @PostMapping("/sayhello")  
    public String sayHello(@RequestBody MyRequest  
request) {  
        return "";  
    }  
}
```

**Effective URL: [http://\[serverip\]/myservice/sayhello](http://[serverip]/myservice/sayhello)**

# Spring Boot Framework

- Uses reflection to first look up all classes with `@RestController` annotation
- Then automatically creates Servlets out of the methods annotated with `@GetMapping`, `@PostMapping`, etc.
- Uses reflection to parse the request parameters into class objects annotated with `@RequestBody`
- Generates the WAR file and launches the Apache Tomcat server
  - Simply execute `mvn spring-boot:run`

```
class MyRequest {  
    private String postDate;  
    private String postContent;  
    // .. Getters and setters  
}  
  
@RestController  
@RequestMapping("/myservice")  
public class MyController {  
    @PostMapping("/sayhello")  
    public String sayHello(@RequestBody MyRequest  
request) {  
        return "";  
    }  
}
```

# Spring Boot demo

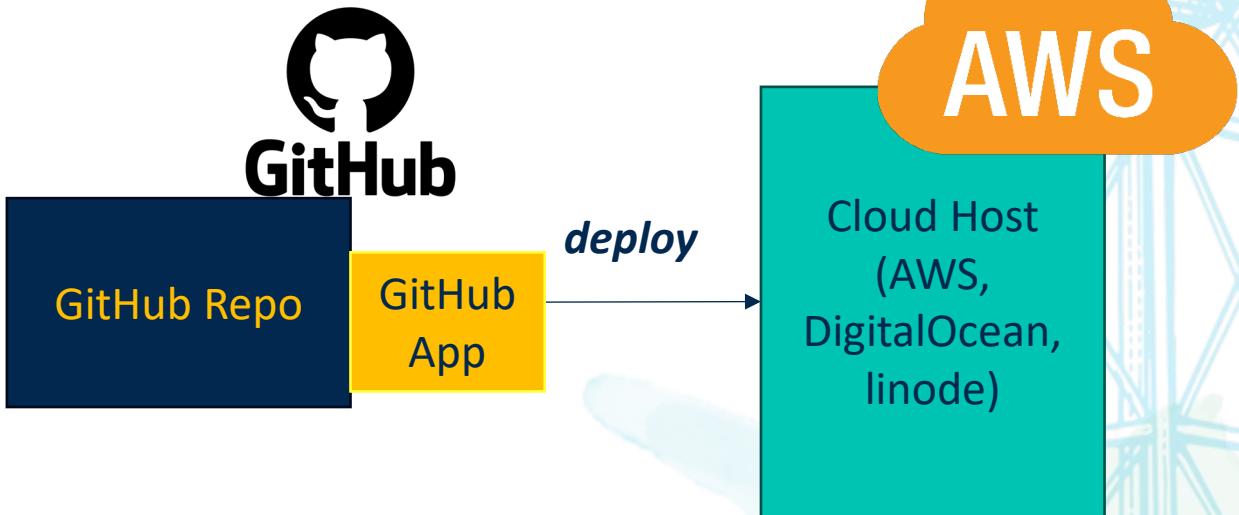
- <https://www.youtube.com/watch?v=uGDhkWc4gYA>

# Continuous Deployment

- Automated release of every code change that passes tests directly to production
- Final stage of a modern CI/CD pipeline
- Uses containers
- GitHub Actions OR GitHub App Integration through cloud host
- GitHub Apps can ***extend the functionality*** of GitHub itself

# Cloud host integration

- Every push to the repo must automatically deploy to the cloud
- How to deploy?
  - Provide a Dockerfile that packages and launches the software
- How to detect new commit pushes?
  - GitHub Actions
  - GitHub Apps provided by the Cloud Host



# Demo – DigitalOcean integration

- DigitalOcean is a cloud provider, like AWS
  - Can launch VMs
- Apps Platform
  - Managed service that automatically builds, deploys, and scales your apps
  - Directly from GitHub
- Following steps similar across many platforms

# Containerize your app

- Containers build on a base image
  - maven:3.9.8-eclipse-temurin-21
    - <https://hub.docker.com/layers/library/maven/3.9.8-eclipse-temurin-21/images/sha256-d86bfd73fbfd7e9f0a9554ae23689bd46289f537df4a5aa368fdd1a8c25dd8d0?context=explore>
  - eclipse-temurin:21-jre

```
# --- Build stage ---
FROM maven:3.9.8-eclipse-temurin-21 AS build
WORKDIR /app
COPY pom.xml . # Copy pom.xml into the container
COPY src ./src
RUN mvn -q -DskipTests package # Build it!

# --- Runtime stage ---
FROM eclipse-temurin:21-jre
WORKDIR /app
COPY --from=build /app/target/*.jar app.jar
CMD ["java", "-jar", "app.jar"]
```

# **Message queues and pub-sub architecture**

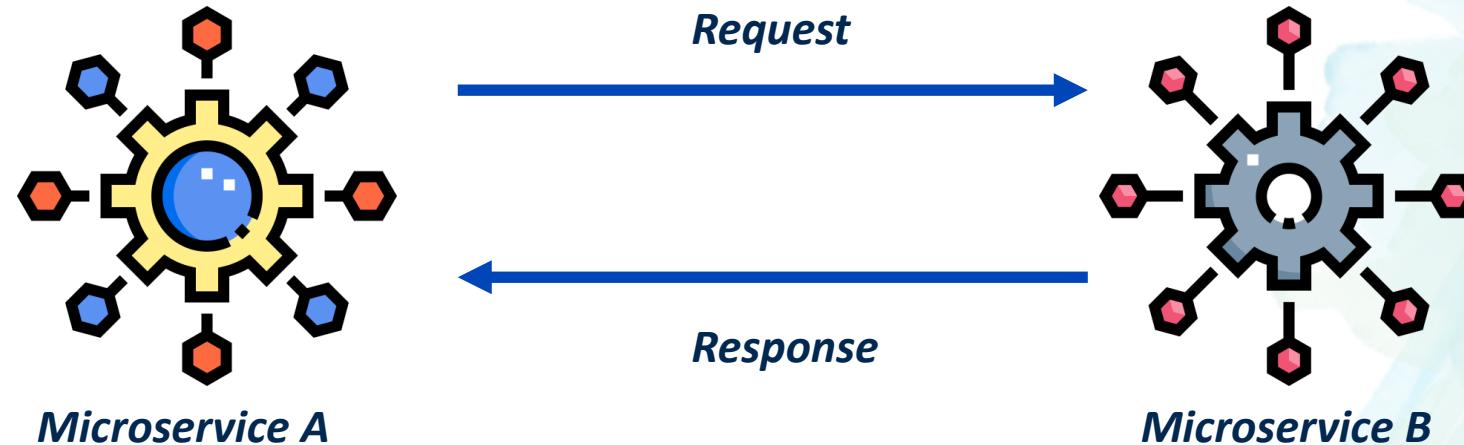
Tapti Palit

# Outline

- RPC
- Message queue
  - Data-structures
  - Persistence
- Pub/sub
  - Data-structures
  - Persistence

# Remote procedure call (RPC)

- Library/framework gives the illusion that the other microservice is running locally
  - Protocol that allows a program to execute a procedure on a remote server as if it were local
- Synchronous communication – caller waits for response before proceeding



# RPC key features

- Language agnostic: the RPC itself does not depend on the service language
- Abstracts network details
  - Typically, over HTTP

# RPC formats

- Text-based (e.g. REST APIs)
  - Uses JSON or XML for data exchange
  - Human-readable, but larger payloads
- Binary formats (e.g. gRPC)
  - Uses binary standards (such as protobuf) for serialization
  - Compact but faster, but less human-readable
  - ... *why?*

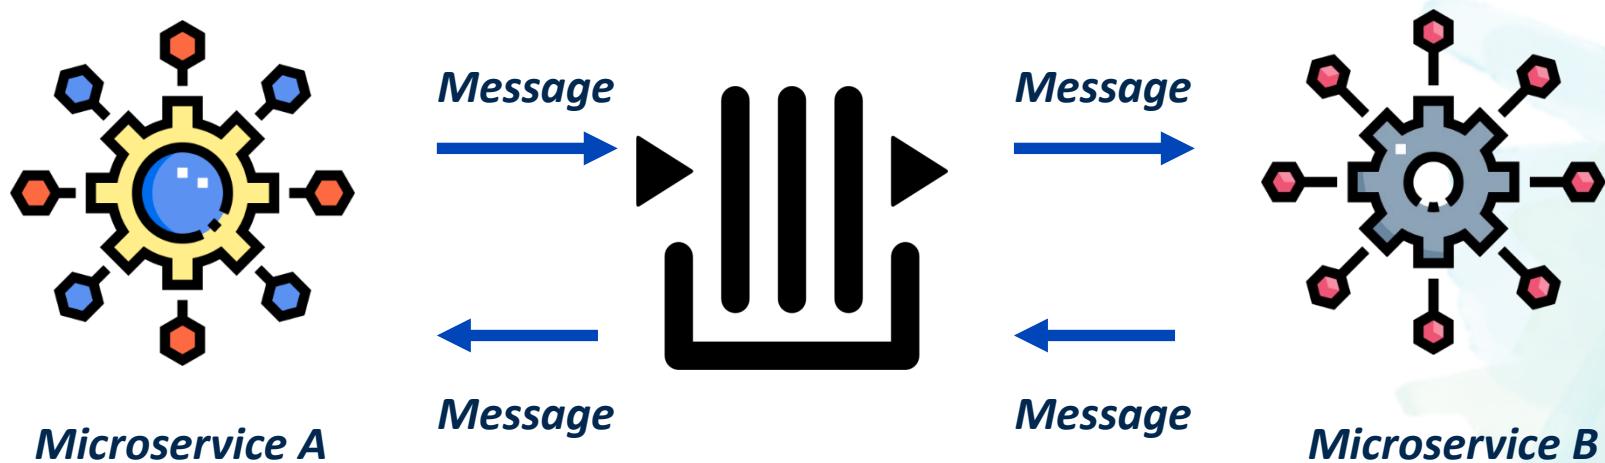
# JSON and REST APIs

- Representational State Transfer (REST) is an architectural style for web services
- Application/microservice exposes an URL
- Uses HTTP methods (GET, POST, PUT, DELETE) to perform operations on resources
- Commonly paired with JSON for data exchange

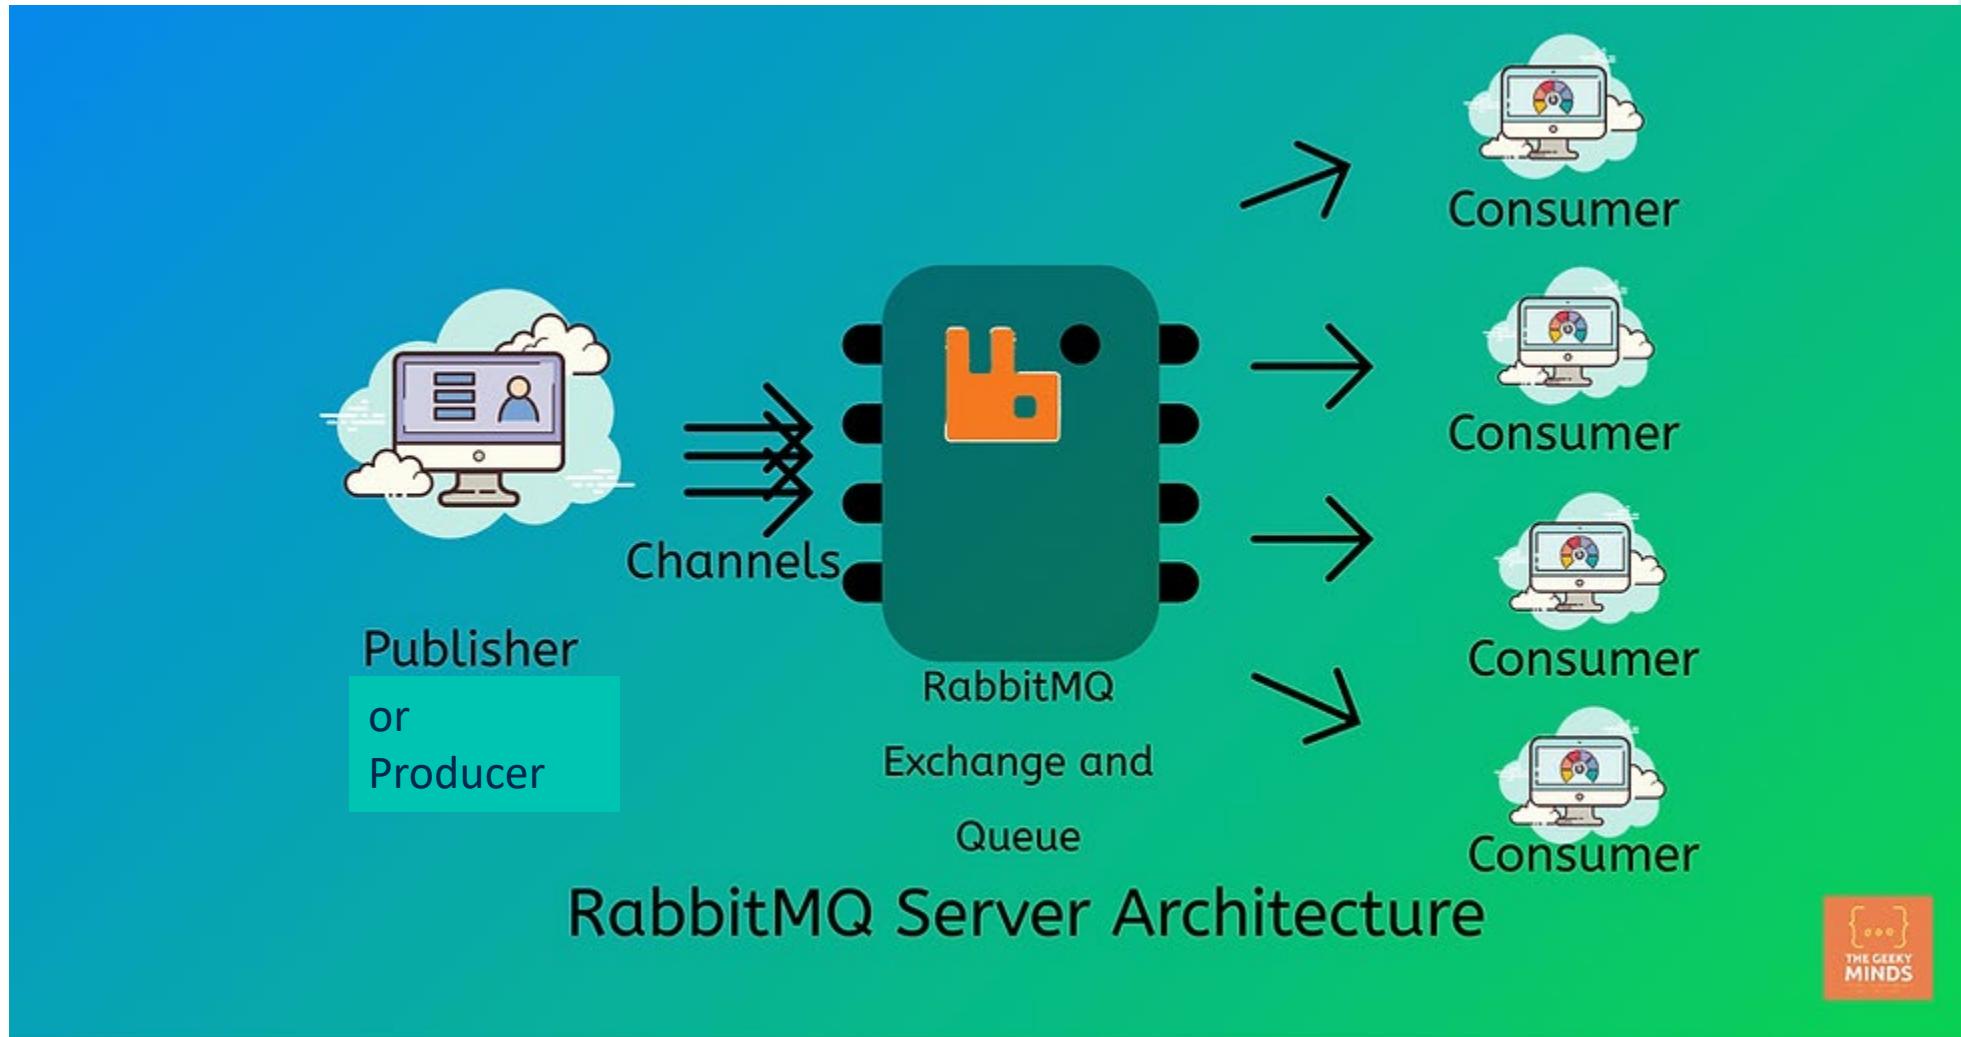
```
// GET Request to Fetch a User denoted by ID  
  
> GET /users/123  
  
// Response  
  
{  
  "id": 123,  
  "name": "John Doe",  
  "email": john.doe@example.com  
}
```

# Message queuing (MQ)

- Asynchronous communication model
  - Messages sent to a queue and processed by consumers independently of the producer
- Stronger decoupling



# RabbitMQ architecture



# RabbitMQ demo

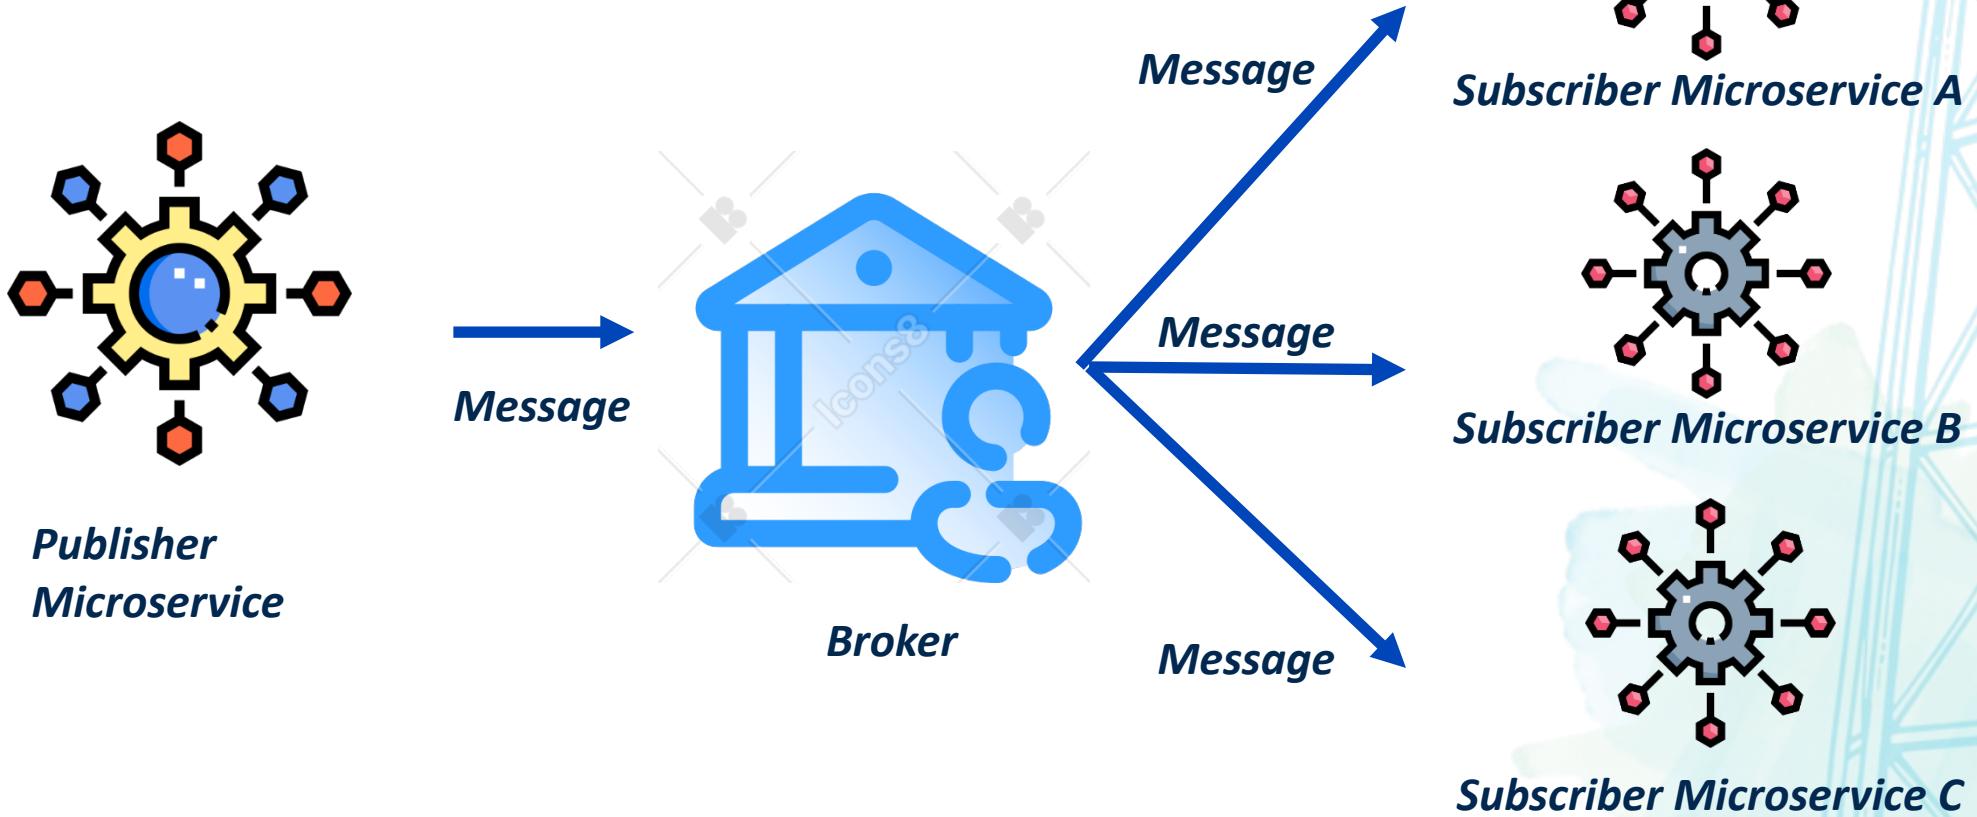
- [https://www.youtube.com/watch?v=WzO6\\_4jeliM](https://www.youtube.com/watch?v=WzO6_4jeliM)



# Pub-sub architecture

- Asynchronous messaging pattern where **publishers** send messages to a central **message broker** or **topic**, and **subscribers** receive messages based on their subscriptions
- Broadcasting: messages can be sent to multiple subscribers
- Typically, messages are persistent at the broker and must be explicitly deleted
- Same frameworks often can act as both MQ or Pub-Sub depending on configuration

# Pub-sub architecture



# Redis pub-sub demo

- [https://www.youtube.com/watch?v=bSe\\_JBSk5w4](https://www.youtube.com/watch?v=bSe_JBSk5w4)



Google Cloud  
Pub/Sub



# Can mix and match!!

- Pipeline architecture using microservices with message queues deployed on a serverless architecture
- Pub-sub architecture with microservices using containers
- Pipeline architecture in one part of the system, with a pub-sub in another
- ... and so on
- Pick what is right for your software system!