

# Modern software architectures

Tapti Palit

# Outline

- Monolithic applications
- Microservices and decentralized data
- Data model and storage engine
  - Relational databases, log-structured merge trees (LSM), event logs (more in Kafka section), in-memory cache
- Communication styles
  - Synchronous, Asynchronous (MQs), Publish-subscribe models


# All about the data!

- Modern software architecture is data-intensive
- Most of the time we'll be concerned about
  - Who owns data?
  - How to optimally store data?
  - How is data shared?
  - How to limit overhead due to data-sharing?


# Monolithic software architecture

- Widely used in early, mid 2000s
- All software components are part of the same application
- All software components run on the same machine
- Typically developed in the same language stack

# Simple online store



# Simple online store



# Relational databases

- Consists of tables
- Each table contains a primary key
  - Database will not allow insertion of two records with same primary key



The background features a faint watermark of a water tower and a bridge.

| Products tbl |             |              |
|--------------|-------------|--------------|
| ProductId    | ProductName | ProductCount |
| 1001         | ABC         | 10           |


| Customer tbl |              |              |
|--------------|--------------|--------------|
| CustomerId   | CustomerName | CustomerAddr |
| 9001         | XYZ          | CDEF         |

| Order tbl |            |           |
|-----------|------------|-----------|
| OrderId   | CustomerId | ProductId |
| 110001    | 9001       | 1001      |

# Foreign keys

- Relational databases maintain relations through foreign keys
- Foreign keys **must refer** to primary keys of other tables
- A table can contain one or more foreign keys

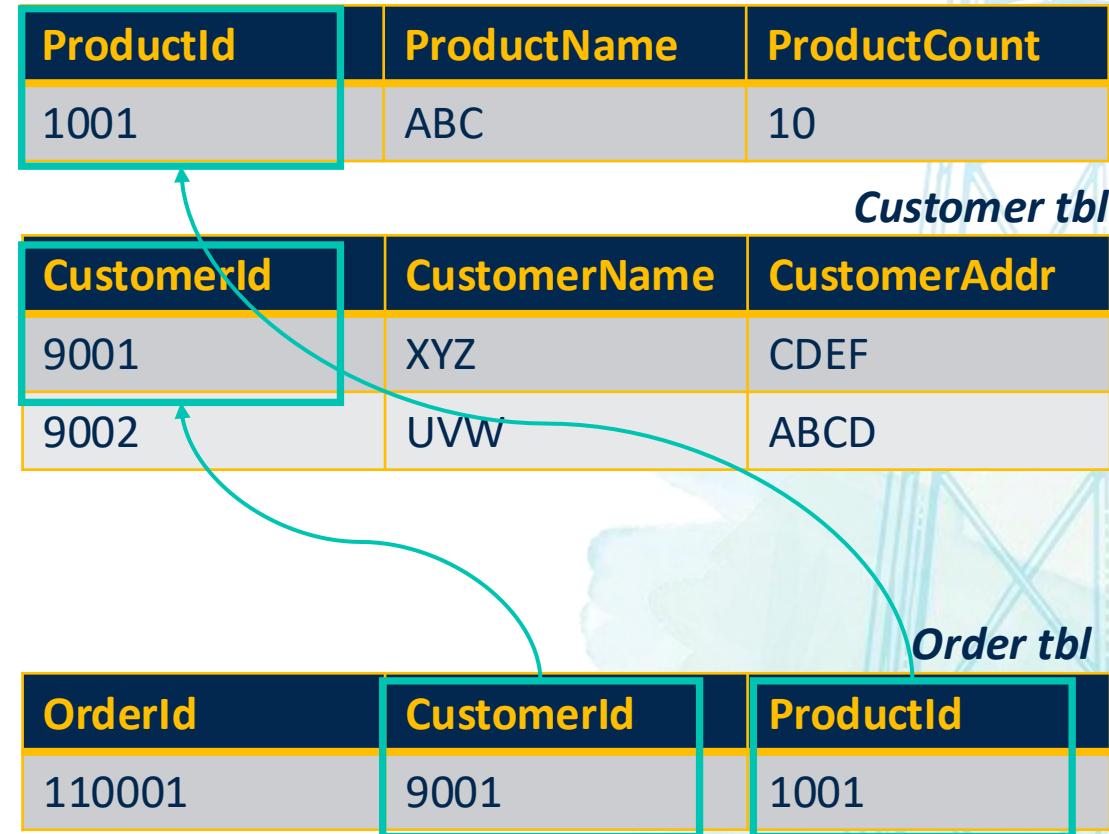


# Structured query language (SQL)

- SQL used to interface between application and database
- ```
CREATE TABLE Products (
    ProductId INT PRIMARY KEY,
    ProductName VARCHAR(100) NOT NULL,
    ProductCount INT NOT NULL);
```
- ```
INSERT INTO Products (ProductId, ProductName,
ProductCount) VALUES (1001, 'ABC', 10);
```

# SQL joins

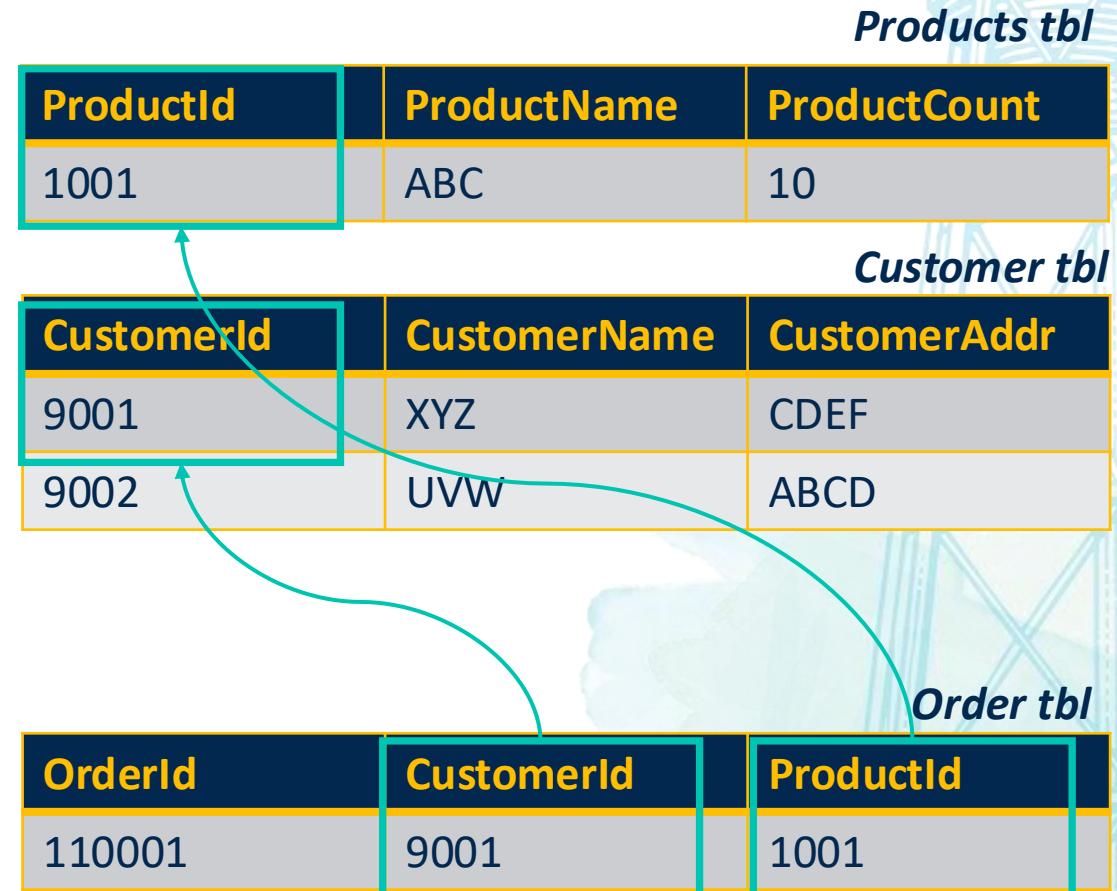
- Allows table joins
- For e.g. find order details for shipping, including product name, customer name, and address



# SQL joins

SELECT

```
    o.OrderId,  
    p.ProductName,  
    c.CustomerName,  
    c.CustomerAddr  
  
FROM Orders o  
JOIN Products p  
    ON o.ProductId = p.ProductId  
JOIN Customers c  
    ON o.CustomerId =  
c.CustomerId;
```



# Need for joins

- Imagine no support for joins
- Order information must contain product name, customer name, and customer address

| ProductId | ProductName | ProductCount |
|-----------|-------------|--------------|
| 1001      | ABC         | 10           |

| CustomerId | CustomerName | CustomerAddr |
|------------|--------------|--------------|
| 9001       | XYZ          | CDEF         |
| 9002       | UVW          | ABCD         |

| OrderId | ProductName | CustomerName | CustomerAddr |
|---------|-------------|--------------|--------------|
| 110001  | ABC         | XYZ          | CDEF         |

# Need for joins

- Lack of join support increases data duplication
- Foreign keys enforce referential integrity
- Data denormalization

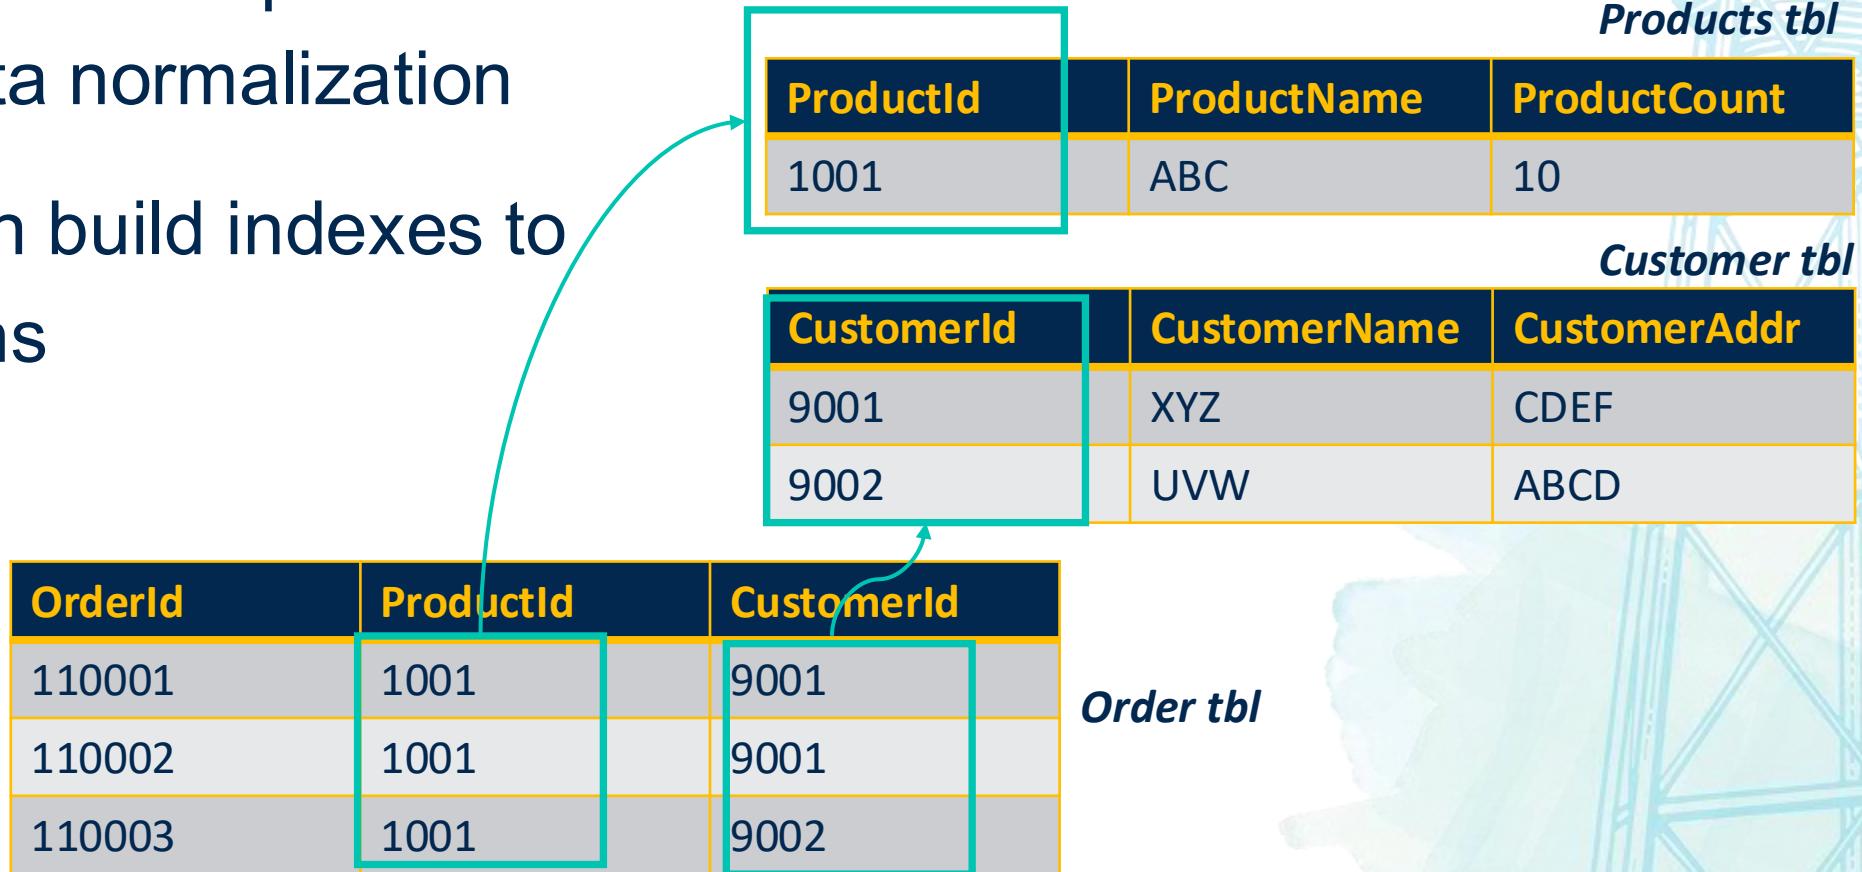
| ProductId | ProductName | ProductCount |
|-----------|-------------|--------------|
| 1001      | ABC         | 10           |

| CustomerId | CustomerName | CustomerAddr |
|------------|--------------|--------------|
| 9001       | XYZ          | CDEF         |
| 9002       | UVW          | ABCD         |

| OrderId | ProductName | CustomerName | CustomerAddr |
|---------|-------------|--------------|--------------|
| 110001  | ABC         | XYZ          | CDEF         |
| 110002  | ABC         | XYZ          | CDEF         |
| 110003  | ABC         | UVW          | ABCD         |

# Normalization


- Joins reduce data duplication and allow data normalization
- Database can build indexes to speed up joins



# SQL != relational databases

- Note: not quite specific to relational databases
- Apache Cassandra, Apache HBase, Apache Kafka + ksqlDB

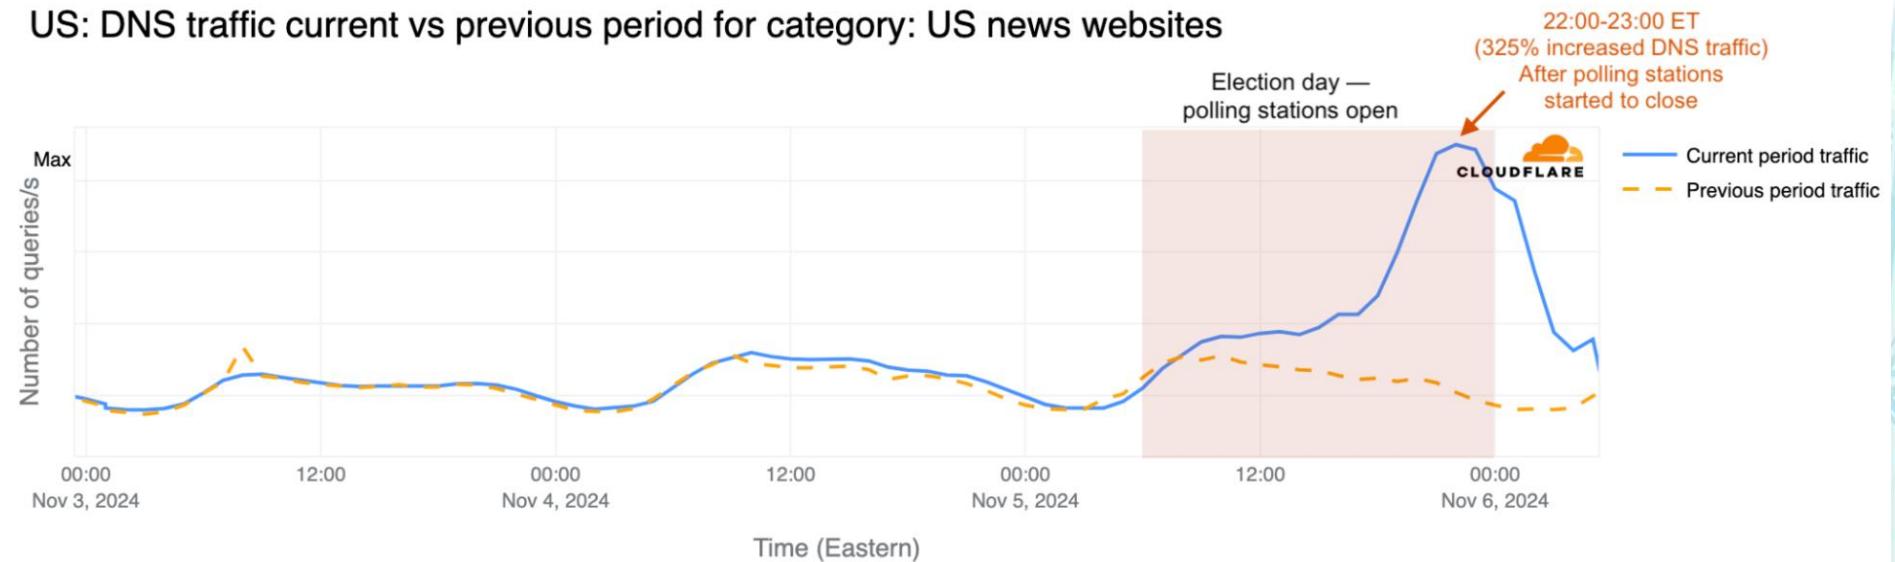
# Simple online store app



# Scalability

- Website traffic is not constant
- Can spike due to planned events
  - Product launches
  - Political events


<https://blog.cloudflare.com/exploring-internet-traffic-shifts-and-cyber-attacks-during-the-2024-us-election/>



# Scalability


- Unplanned events
  - Blog post goes viral
  - System architecture should *adapt* to handle such events

<https://www.residualthoughts.com/2018/05/20/traffic-data-from-a-viral-post/>




# Traffic spike to online store

*At high utilizations, queuing delays spike tail latency  
Many workloads aim to operate at 40-50% utilizations*




# Vertical scaling

- Use more powerful machines
- Add more CPUs, DRAM, network bandwidth



# Vertical scaling limitations

Resource wastage (assuming we aim for 40% usage)




# Vertical scaling limitations

- Twitter has 200M-300M active users per day
- No single machine, no matter how powerful, can support that
- Goal: autoscaling
  - **Dynamically** spawn new machines during **high loads**
  - Not possible using vertical scaling alone (modulo virtual machines, containers)
  - More in Kubernetes module

# Horizontal scaling for monolithic apps

Add more machines and replicate application on each machine




# Load balancer

- Distributes incoming requests across multiple servers to improve scalability and availability
- Strategies
  - Round Robin – Sequentially routes requests across servers; simple but doesn't account for server load
  - Least Connections – Directs traffic to the server with the fewest active connections; adapts well to uneven load
  - Least Response Time – Chooses the server with the fastest response time and fewest connections; performance-oriented
  - Random Policy – Selects servers randomly; useful in stateless, uniform environments
  - Weighted Distribution – Allocates requests based on server capacity (e.g., CPU power, memory)
- Each strategy has tradeoffs
- More details in Kubernetes module


# Autoscaling

Auto-scale to more machines during traffic spike (more in Kubernetes module)



# Autoscaling

Minimum resource wastage



*Did we solve all problems?*

# Heterogenous resource requirements across services

- Provisioning is driven by the most resource-hungry service
- Example RAM requirements
  - CustomerService: 32 GB
  - OrderService: 18 GB
  - ProductService: 16 GB
  - Minimum machine RAM? 32 GB

# Deployability concerns

- Updating one component requires redeploying the entire application
- Reverting a change requires redeploying the entire application
- Slow, error-prone process

# Need for low interdependence

- Software often consists of thousands of components
  - Each component has a dedicated team working on it
- Teams need to work independently
  - ProductService team should be able to update the Products Tbl schema without consulting
- Need low coupling between services
- Solution: microservices

# Microservices

- An approach to developing a single application as a suite of small services, each running in its own process and communicating with lightweight mechanisms, often an HTTP resource API
- Independently deployable by automated processes
- Bare minimum centralized management
- Smart endpoints connected by “dumb” pipes

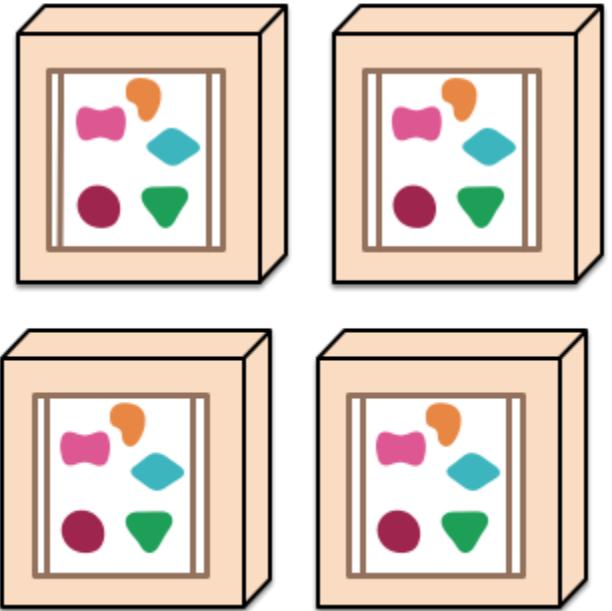
<https://martinfowler.com/articles/microservices.html>

# Microservices overview


*A monolithic application puts all its functionality into a single process...*




*... and scales by replicating the monolith on multiple servers*



*A microservices architecture puts each element of functionality into a separate service...*




*... and scales by distributing these services across servers, replicating as needed.*




<https://martinfowler.com/articles/microservices.html>

# Microservice architecture for online store




# Each microservice chooses its own database




# What did we achieve?

- Decompose monolithic app into microservices
- Improve decoupling
  - Each microservice can be scaled independently
  - Each microservice can be deployed independently
  - Each microservice can evolve independently – DB schema, choice of programming languages
- ***Did we solve all problems?***

# Shipping an order



# Shipping an order



*Cannot perform joins!*


# Solution – denormalization of data

- Must denormalize the data
- DO NOT want to query the Customer and Product microservices when shipping
- Will increase network overhead

| OrderId | ProductName | CustomerName | CustomerAddr |
|---------|-------------|--------------|--------------|
| 110001  | ABC         | XYZ          | CDEF         |
| 110002  | ABC         | XYZ          | CDEF         |
| 110003  | ABC         | UVW          | ABCD         |

# Solution – denormalization of data

Typical solution: user confirms the Customer data on the Order HTML page




The screenshot shows an 'Order Details' page from the Amazon website. At the top, there's a navigation bar with links like 'All', 'Kindle Books', 'Epic Daily Deals', 'Smart Home', 'Amazon Launchpad', 'Prime', 'Prime Video', 'Books', 'Pet Supplies', 'Groceries', 'Handmade', 'Livestreams', and 'Shop Epic Daily Deals now'. Below the navigation, it says 'Your Account > Your Orders > Order Details'. The main section is titled 'Order Details' and shows an order placed on October 5, 2021, with Order # [redacted]. It includes sections for 'Shipping Address' (Dave Johnson, United States), 'Payment Method' (partially visible), and 'Order Summary' (Subtotal, Shipping & Handling, Total before tax, Estimated tax to be collected). A red arrow points from the 'CustomerName' column in the adjacent table to the 'CustomerName' field in the 'Shipping Address' section of the screenshot.

| OrderId | ProductName | CustomerName | CustomerAddr |
|---------|-------------|--------------|--------------|
| 110001  | ABC         | XYZ          | CDEF         |
| 110002  | ABC         | XYZ          | CDEF         |
| 110003  | ABC         | UVW          | ABCD         |

# Consistency guarantees

- Consistency property of a system governs how and when updates to shared data become visible to different components of the system
- Monolithic apps typically have strong consistency – updates reflect immediately




| ProductId | ProductName | ProductCount |
|-----------|-------------|--------------|
| 1001      | ABC         | 10           |

| CustomerId | CustomerName | CustomerAddr |
|------------|--------------|--------------|
| 9001       | XYZ          | CDEF         |
| 9002       | UVW          | ABCD         |

| OrderId | ProductId | CustomerId |
|---------|-----------|------------|
|         |           |            |

# Consistency guarantees

- Consistency property of a system governs how and when updates to shared data become visible to different components of the system
- Monolithic apps typically have strong consistency – updates reflect immediately



| ProductId | ProductName | ProductCount |
|-----------|-------------|--------------|
| 1001      | ABC         | 10           |

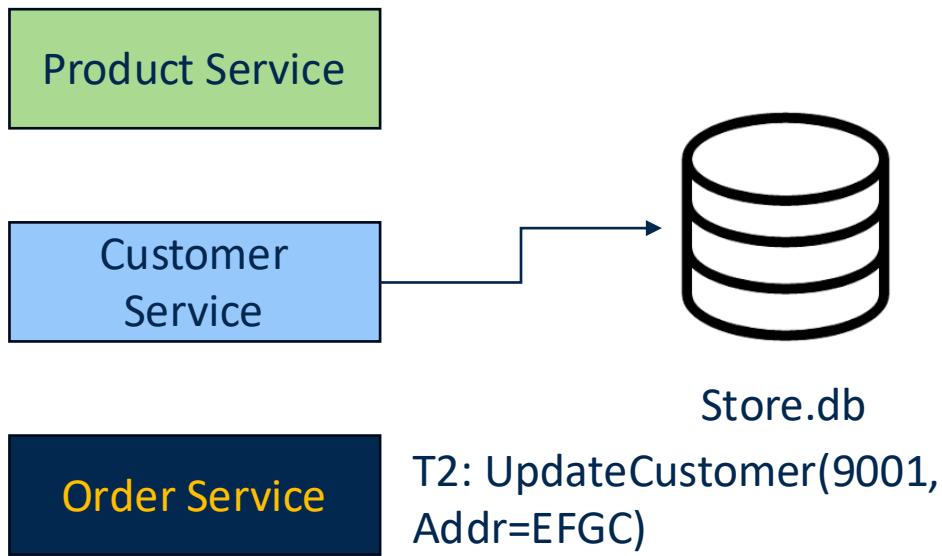
| CustomerId | CustomerName | CustomerAddr |
|------------|--------------|--------------|
| 9001       | XYZ          | CDEF         |
| 9002       | UVW          | ABCD         |

| OrderId | ProductId | CustomerId |
|---------|-----------|------------|
| 110001  | 1001      | 9001       |

# Consistency guarantees

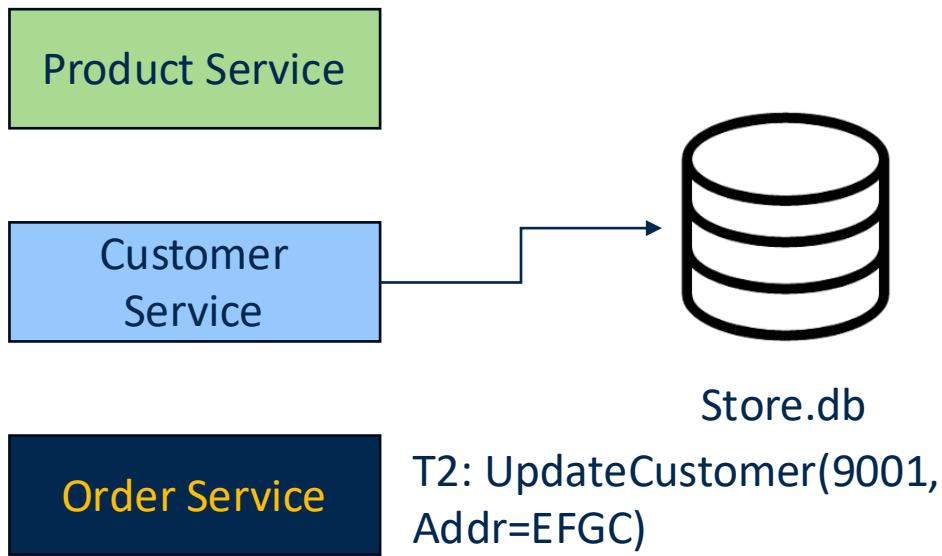
- Consistency property of a system governs how and when updates to shared data become visible to different components of the system
- Monolithic apps typically have strong consistency – updates reflect immediately



| ProductId  | ProductName  | ProductCount |
|------------|--------------|--------------|
| 1001       | ABC          | 10           |
| CustomerId | CustomerName | CustomerAddr |
| 9001       | XYZ          | CDEF         |
| 9002       | UVW          | ABCD         |
| OrderId    | ProductId    | CustomerId   |
| 110001     | 1001         | 9001         |

# Consistency guarantees


- Consistency property of a system governs how and when updates to shared data become visible to different components of the system
- Monolithic apps typically have strong consistency – updates reflect immediately



| ProductId  | ProductName  | ProductCount |
|------------|--------------|--------------|
| 1001       | ABC          | 10           |
| CustomerId | CustomerName | CustomerAddr |
| 9001       | XYZ          | CDEF         |
| 9002       | UVW          | EFGC         |
| OrderId    | ProductId    | CustomerId   |
| 110001     | 1001         | 9001         |

# Consistency guarantees

- Consistency property of a system governs how and when updates to shared data become visible to different components of the system
- Monolithic apps typically have strong consistency – updates reflect immediately to subsequent reads



T3: ShipOrder ->  
JOIN(1001, 9001) ->  
(ABC, UVW, EFGC)

| ProductId | ProductName | ProductCount |
|-----------|-------------|--------------|
| 1001      | ABC         | 10           |


| CustomerId | CustomerName | CustomerAddr |
|------------|--------------|--------------|
| 9001       | XYZ          | CDEF         |
| 9002       | UVW          | EFGC         |

| OrderId | ProductId | CustomerId |
|---------|-----------|------------|
| 110001  | 1001      | 9001       |


# Microservice consistency guarantees

- Microservice-oriented apps are distributed
- Enforcing strong consistency guarantees in distributed systems is very hard due to network overhead, network partition, node failures
- Typically have eventual consistency guarantees – updates are eventually visible to all components
- System must work around these limitations
- Core idea: a microservice **owns** some data
  - **Notifies** dependent microservices of change
  - **No guarantees** on when the updates synchronize


# Working with eventual consistency




# Working with eventual consistency



# Working with eventual consistency



# Working with eventual consistency




# Address confirmation page


**Name: UVW,  
Address: ABCD**

The screenshot shows the Amazon Order Details page. At the top, the navigation bar includes the Amazon Prime logo, a search bar, and links for Hello, Dave, Account & Lists, Returns & Orders, and a Cart with 0 items. Below the navigation, the breadcrumb trail reads Your Account > Your Orders > Order Details. The main section is titled "Order Details" and shows the order was placed on October 5, 2021, with Order# [REDACTED]. A "View or Print invoice" button is available. The "Shipping Address" section displays "Dave Johnson" and "United States". A red arrow points from the text "Address: ABCD" in the previous slide to the "Change" link next to the United States dropdown. The "Payment Method" section shows a placeholder for a credit card. The "Order Summary" section provides a breakdown of costs: Item(s) Subtotal, Shipping & Handling, Total before tax, Estimated tax to be collected, and Grand Total. Below the summary, a product item is listed: "Arriving tomorrow by 10pm" - "Amazon Basics Neoprene Coated Dumbbell Hand Weight Set, 10-Pound, Set of 2, Navy Blue". It includes details like Sold by Amazon.com Services, LLC, \$34.99, Condition: New, and an "Add gift option" link. A "Buy it again" button is also present. To the right of the product, there is a sidebar with buttons for "Track package", "Change Payment Method", "Change shipping speed", "Cancel items", and "Archive order".


# Working with eventual consistency




# Working with eventual consistency




# Working with eventual consistency



# Working with eventual consistency



# Working with eventual consistency




# Microservice design concerns


- Solution is generally use-case dependent
- General goal - limit network communication

# Network communication cost

Monolithic applications



Microservices




# Network overhead decomposition

- Network latency
  - Datacenters use fast network connections
    - Infiniband has ~1-2 microsecond latency, 400+ Gbps bandwidth
    - Still not as fast as a local function call
  - Kernel overhead

# Kernel overhead for network comm.

- The kernel contains the networking code
- TCP-IP is the most common networking stack
- It is organized in layers
- Every layer has a **protocol**



# Each layer has a protocol

- Every layer/protocol has a fixed message format

- Header


- Payload

- [optional] Checksum




- As the packet traverses through the layers, packets are rewapped

# Life of a packet



# Life of a packet



*Packet encapsulation process can take  
100+ microseconds*

*Reading 6 will discuss an alternative*



# Microservice pros

- Stronger decoupling and lower interdependence
- Improved scalability
- Easier deployment


# Microservice cons

- Causes data denormalization
- Network overhead
- Higher complexity
- Debugging complex interactions is harder

# Latency vs throughput

- Latency – time taken for one operation
  - Measured in seconds, milliseconds, microseconds, etc
- Throughput – number of operations performed in unit time
  - Rps – requests per second
  - Ops – operations per second
- Microservices increase latency compared to monolithic operation
- Microservices significantly improve throughput
- Note: throughput of a system is max when utilization  $\rightarrow 100\%$  (memory, CPU, network), but latency spikes dramatically due to ***queuing effect***


# Microservices at Uber (2019)



<https://x.com/msuriar/status/1110244877424578560>

# Microservices at Amazon (2008)

- Code-named “Deathstar”



<https://x.com/Werner/status/741673514567143424>

# Data management and communication

- Once data and computation are split across services, two problems remain
  - How is data stored?
  - How services communicate?

# Spring Boot overview

# HTTP GET and POST request

- GET request - Used to retrieve data from the server

- GET /index.html HTTP/1.1

*URL*

- GET /index.html?name=ECS160&data=02132025

*Request parameters*

# HTTP GET and POST request

- POST request - used to send data to the server

**URL**

- POST `/users` HTTP/1.1

Content-Type: application/json{

```
"name": "John Doe",  
"email": john.doe@example.com  
}
```

***Post “body”***

# Spring Boot Overview

- Framework for creating RESTful microservices
- Reduces boilerplate configuration code
- Embedded server (Tomcat/Jetty)
- Simplifies microservice creation through annotations
- Built-in support for REST APIs

# RESTful microservices with Spring Boot

- Create classes that can act as REST endpoints
- Uses annotations to denote REST endpoint URLs
  - Allows complete decoupling from the boilerplate code
- Types of requests
  - @GetMapping, @PostMapping, @PutMapping, and so on... for all HTTP methods
- @PathVariable – extract variable from GET request
- @RequestBody – extract the post request body

```
class MyRequest {  
    private String postDate;  
    private String postContent;  
    // ... Getters and setters  
}  
  
@RestController  
@RequestMapping("/myservice")  
public class MyController {  
    @PostMapping("/sayhello")  
    public String sayHello(@RequestBody MyRequest  
request) {  
        return "";  
    }  
}
```

**Effective URL: [http://\[serverip\]/myservice/sayhello](http://[serverip]/myservice/sayhello)**

# Spring Boot Framework


- Uses reflection to first look up all classes with `@RestController` annotation
- Then automatically creates Servlets out of the methods annotated with `@GetMapping`, `@PostMapping`, etc.
- Uses reflection to parse the request parameters into class objects annotated with `@RequestBody`
- Generates the WAR file and launches the Apache Tomcat server
  - Simply execute `mvn spring-boot:run`

```
class MyRequest {  
    private String postDate;  
    private String postContent;  
    // ... Getters and setters  
}  
  
@RestController  
@RequestMapping("/mvservice")  
public class MyController {  
    @PostMapping("/sayhello")  
    public String sayHello(@RequestBody MyRequest  
request) {  
        return "";  
    }  
}
```

# Software communication styles

# Remote procedure call (RPC)

- Library/framework gives the illusion that the other microservice is running locally
  - Protocol that allows a program to execute a procedure on a remote server as if it were local
- Synchronous communication – caller waits for response before proceeding



# RPC key features

- Language agnostic: the RPC itself does not depend on the service language
- Abstracts network details
  - Typically, over HTTP

# RPC formats

- Text-based (e.g. REST APIs)
  - Uses JSON or XML for data exchange
  - Human-readable, but larger payloads
- Binary formats (e.g. gRPC)
  - Uses binary standards (such as protobuf) for serialization
  - Compact but faster, but less human-readable
    - ... *why?*


# JSON and REST APIs

- Representational State Transfer (REST) is an architectural style for web services
- Application/microservice exposes an URL
- Uses HTTP methods (GET, POST, PUT, DELETE) to perform operations on resources
- Commonly paired with JSON for data exchange

```
// GET Request to Fetch a User denoted by ID  
  
> GET /users/123  
  
// Response  
  
{  
    "id": 123,  
    "name": "John Doe",  
    "email": john.doe@example.com  
}
```

# Message queuing (MQ)

- Asynchronous communication model
  - Messages sent to a queue and processed by consumers independently of the producer
- Stronger decoupling




# Open source and paid MQ services



# Pub-sub architecture

- Asynchronous messaging pattern where **publishers** send messages to a central **message broker** or **topic**, and **subscribers** receive messages based on their subscriptions
- Broadcasting: messages can be sent to multiple subscribers
- Typically, messages are persistent at the broker and must be explicitly deleted
- Same frameworks often can act as both MQ or Pub-Sub depending on configuration


# Pub-sub architecture



# Pub-sub frameworks



# RabbitMQ architecture



# Asynchronous communication with RabbitMQ

```
public class RabbitMQProducer {  
    private static final String QUEUE_NAME = "hello_queue";  
  
    public static void main(String[] args) {  
        // Create a connection factory  
        ConnectionFactory factory = new ConnectionFactory();  
        factory.setHost("localhost"); // RabbitMQ server  
        address  
  
        try (Connection connection = factory.newConnection();  
             Channel channel = connection.createChannel()) {  
  
            // Declare a queue (it must exist before publishing)  
            channel.queueDeclare(QUEUE_NAME, false, false,  
            false, null);  
  
            String message = "Hello, RabbitMQ!";  
            channel.basicPublish("", QUEUE_NAME, null,  
            message.getBytes());  
  
            System.out.println(" [x] Sent: '" + message + "'");  
        } catch (Exception e) {  
            e.printStackTrace();  
        }  
    }  
}
```

```
public class RabbitMQConsumer {  
    private static final String QUEUE_NAME = "hello_queue";  
  
    public static void main(String[] args) {  
        ConnectionFactory factory = new ConnectionFactory();  
        factory.setHost("localhost");  
  
        try {  
            Connection connection = factory.newConnection();  
            Channel channel = connection.createChannel();  
  
            // Declare the queue in case it doesn't exist  
            channel.queueDeclare(QUEUE_NAME, false, false, false,  
            null);  
  
            System.out.println(" [*] Waiting for messages...");  
  
            // Create a consumer to receive messages  
            DeliverCallback deliverCallback = (consumerTag, delivery) -> {  
                String message = new String(delivery.getBody(), "UTF-  
                8");  
  
                System.out.println(" [v] Received: '" + message + "'");  
            };  
  
            channel.basicConsume(QUEUE_NAME, true, deliverCallback,  
            consumerTag -> { });  
  
        } catch (Exception e) {  
            e.printStackTrace();  
        }  
    }  
}
```

# Can mix and match!!

- Pipeline architecture using microservices with message queues deployed on a serverless architecture
- Pub-sub architecture with microservices using containers
- Pipeline architecture in one part of the system, with a pub-sub in another
- ... and so on
- Pick what is right for your software system!