Software Testing

Tapti Palit

Software has bugs

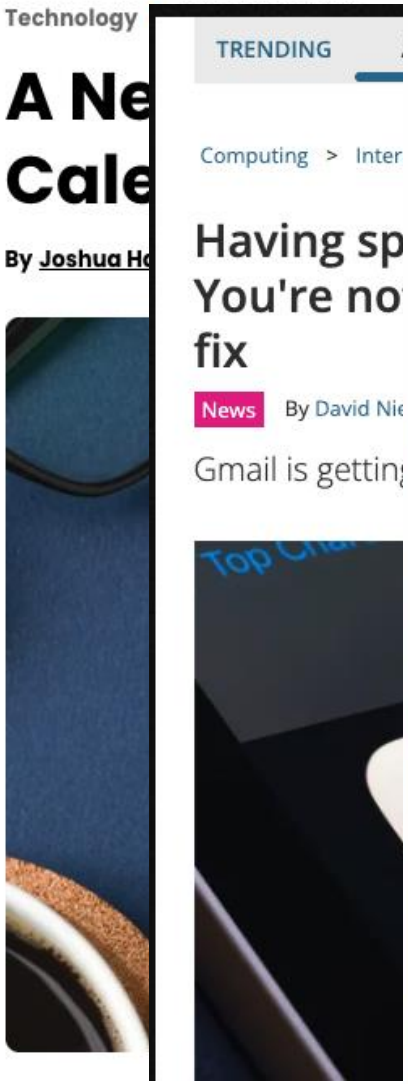

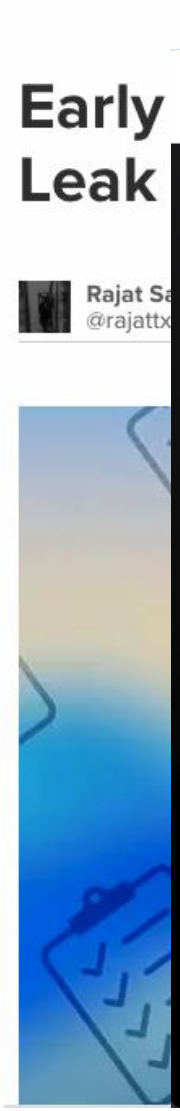

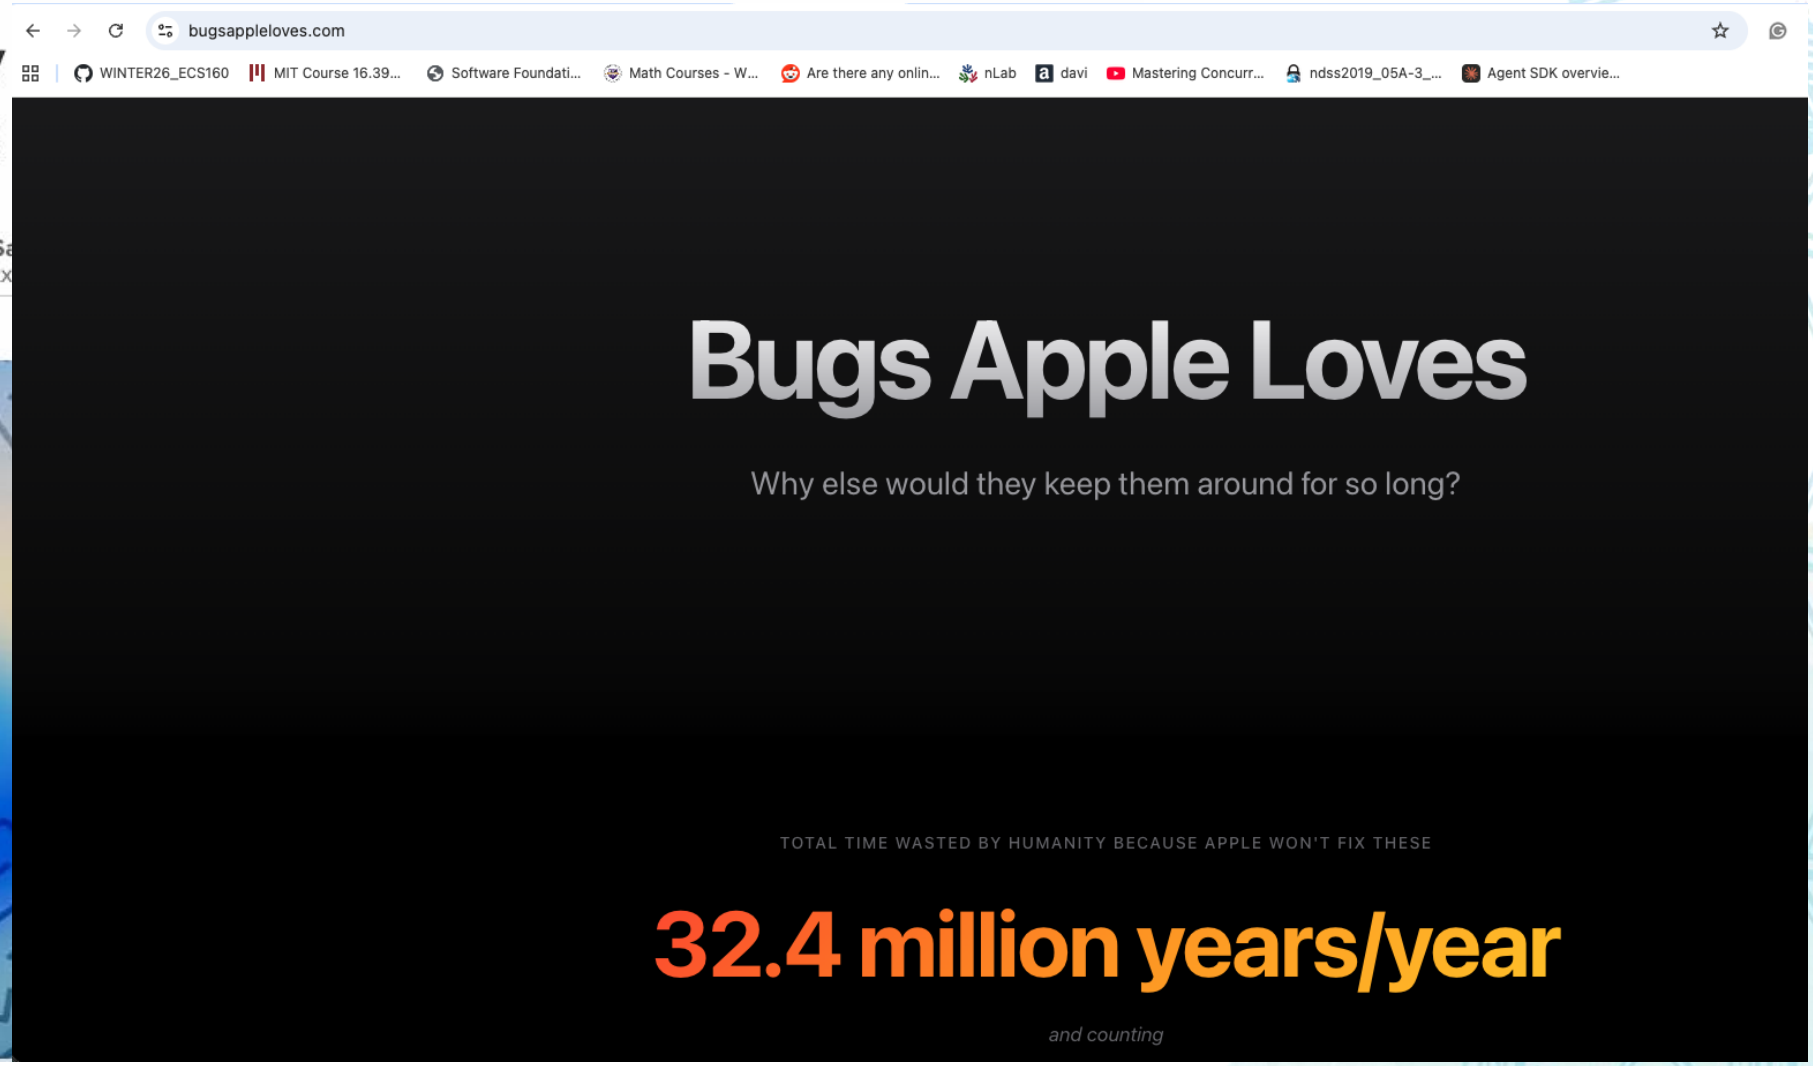

Some bugs are security vulnerabilities

cybernews.com/security/high-severity

26_ECS160 MIT Course 16.39... Soft

Publi

Image

NGII
enab
redir

Malware

NEWS

Upda
day b
execu
webp

by Pieter Ar

Google has
in how the

CVE-2026-
it serious e
release.

Home > News > Microsoft > Microsoft patches actively exploited Office zero-day vulnerability

Micro

By Sergiu

Microsoft has released emergency out-of-band security updates to patch a high-severity Microsoft Office

CYBERCRIME MAGAZINE

ABOUT RESEARCH NEWS

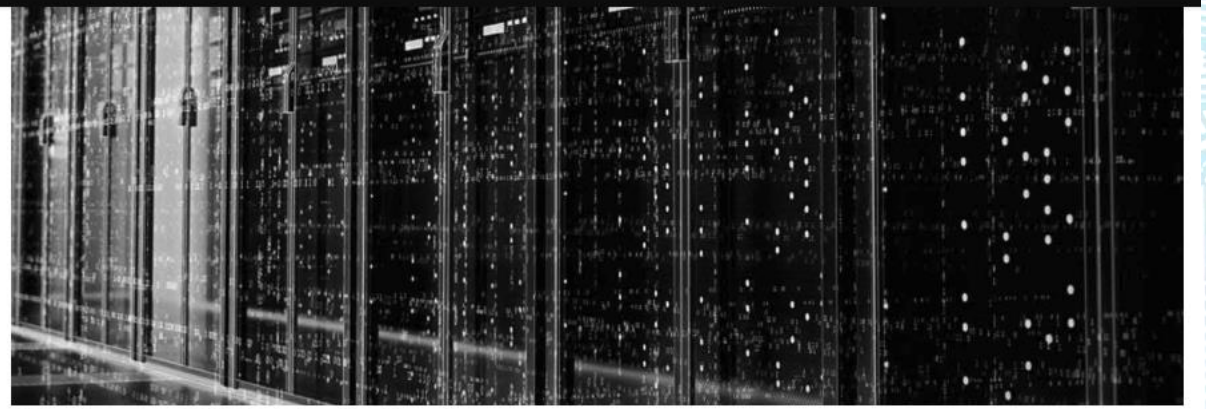

Cybercrime Costs. PHOTO: Cybercrime Magazine.

Cybercrime To Cost The World $10.5 Trillion Annually By 2025

Special Report: Cyberwarfare In The C-Suite.

– *Steve Morgan, Editor-in-Chief*

Sausalito, Calif. – Nov. 13, 2020

If it were measured as a country, then cybercrime — which is predicted to inflict damages totaling $6 trillion USD globally in 2021 — would be the world's third-largest

Software testing

- Process of finding bugs and vulnerabilities in software
- *"Program testing can be used to show the presence of bugs, but never to show their absence"*
  - By E. W. Dijkstra

Types of software testing

- Testing based on dynamic analysis
  - Software is executed
  - Monitor for assertion failures and other silent bugs during execution
  - Unit testing (JUnit), integration testing, ***fuzz testing***
- Testing based on static analysis
  - Source code is analyzed without execution
  - Analysis identifies known bug patterns (un-itialized variables, unsafe data flow, null pointer dereferences)
  - ***Dataflow analysis, control flow analysis, abstract interpretation***
- Hybrid approach: ***symbolic execution***

Dynamic Testing

Recap: unit testing with JUnit

- Test individual functions
- Add assertions to check if post-conditions hold
- Test corner-cases
- Manually created

```
import org.junit.jupiter.api.Test;
import
org.junit.jupiter.api.Assertions.assertEquals;

public class CarTest {
    @Test
    public void testAccelerate() {
        Car car = new Car("Toyota", 50);
        car.accelerate();
        assertEquals(car.getSpeed(), 60);
    }
}
```

Hard to cover all code with unit testing

- Linux kernel
  - ~30 million lines of code
- Google's chromium browser
  - ~28 million lines of code
- Hard to manually write test-cases to cover all this code
- ***Need automation***

Detour: silent failures

- Some bugs lead to crashes, assertion failures
  - Detecting crash bugs requires exercising the buggy code with the right inputs
  - E.g, `assert(i > 0)`
- Some bugs are silent

Detour: silent failures

C/C++

```
#include <stdio.h>

int main(int argc, char** argv) {
    char a = ‘A’;
    char arr[4] = {‘a’, ‘b’, ‘c’,
‘d’};
    int index = atoi(argv[1]);
    printf(“%c\n”, arr[index]);

}

> gcc test.c –o test
> ./test 10
> Output???
```

Java

```
class MyApp {
    public static void main(String[] args) {
        char a = ‘A’;
        char arr[] = {‘a’, ‘b’, ‘c’, ‘d’};
        int index = Integer.parseInt(args[0]);
        System.out.println(arr[index]);
    }
}

> javac MyApp.java
> java MyApp.class 10
> # Output?
```

Python

```
import sys
arr = [‘a’, ‘b’, ‘c’, ‘d’]
index = int(sys.argv[1])
print(arr[index])
> python3 test.py 10
```

Detour: silent failures

C/C++

```
#include <stdio.h>

int main(int argc, char** argv) {
    char a = 'A';
    char arr[4] = {'a', 'b', 'c', 'd'};
    int index = atoi(argv[1]);
    printf("%c\n", arr[index]);

}

> gcc test.c -o test
> ./test 10
>
```

Java

```
class MyApp {
    public static void main(String[] args) {
        char a = 'A';
        char arr[] = {'a', 'b', 'c', 'd'};
        int index = Integer.parseInt(args[0]);
        System.out.println(arr[a]);
    }
}

> javac MyApp.java
> java MyApp.class 10
> Exception in thread "main"
java.lang.ArrayIndexOutOfBoundsException: Index 10 out of bounds
for length 4
```

Undefined Behavior (UB)

Python

```
import sys
arr = ['a', 'b', 'c', 'd']
index = int(sys.argv[1])
print(arr[index])
> python3 test.py 10
IndexError: list index out of range
```

Detour: silent failures

- Some bugs lead to crashes, assertion failures
  - Detecting crash bugs requires exercising the buggy code with the right inputs
  - E.g, `assert(i > 0)`
- Some bugs are silent
  - Detecting silent (UB) bugs requires both, exercising the buggy code with the right inputs AND an oracle to identify the buggy execution
  - E.g., `arr[index]` requires `check(index < length(arr))`

Fuzz testing

- Goal: to find application inputs that reveal bugs
- Test the entire application instead of individual functions
- Generate inputs randomly until an error is uncovered

Fuzz testing

Source for random input on *nix systems

Black-box testing

Types of bugs

- Incorrect argument validation, incorrect type casting
- Buffer overflows, memory leak, div-by-zero, use-after-free, segmentation fault, and so on…
- Logic errors through assertion failures
- ***Applications written in any language can be fuzzed***

Pros and cons

- Pros: generating random inputs is simple
- Cons: random inputs are not good inputs
- E.g., fuzzing a SQL database software
  - Real inputs: `CREATE TABLE mytable AS (id INTEGER, … );`
  - Random inputs: `aVkjW3txpLZ044zo5kLdUlsU2MtlLNkhwxI8Aew7c0KPfTS115i2rzfBlgod`
  - The database will reject such statements as invalid SQL and not execute the SQL command -> the SQL code will not even be tested

Mutation fuzzer - generate inputs via mutation

- Start with some known valid inputs called *"seeds"*
- Keep mutating them based on some "strategy"

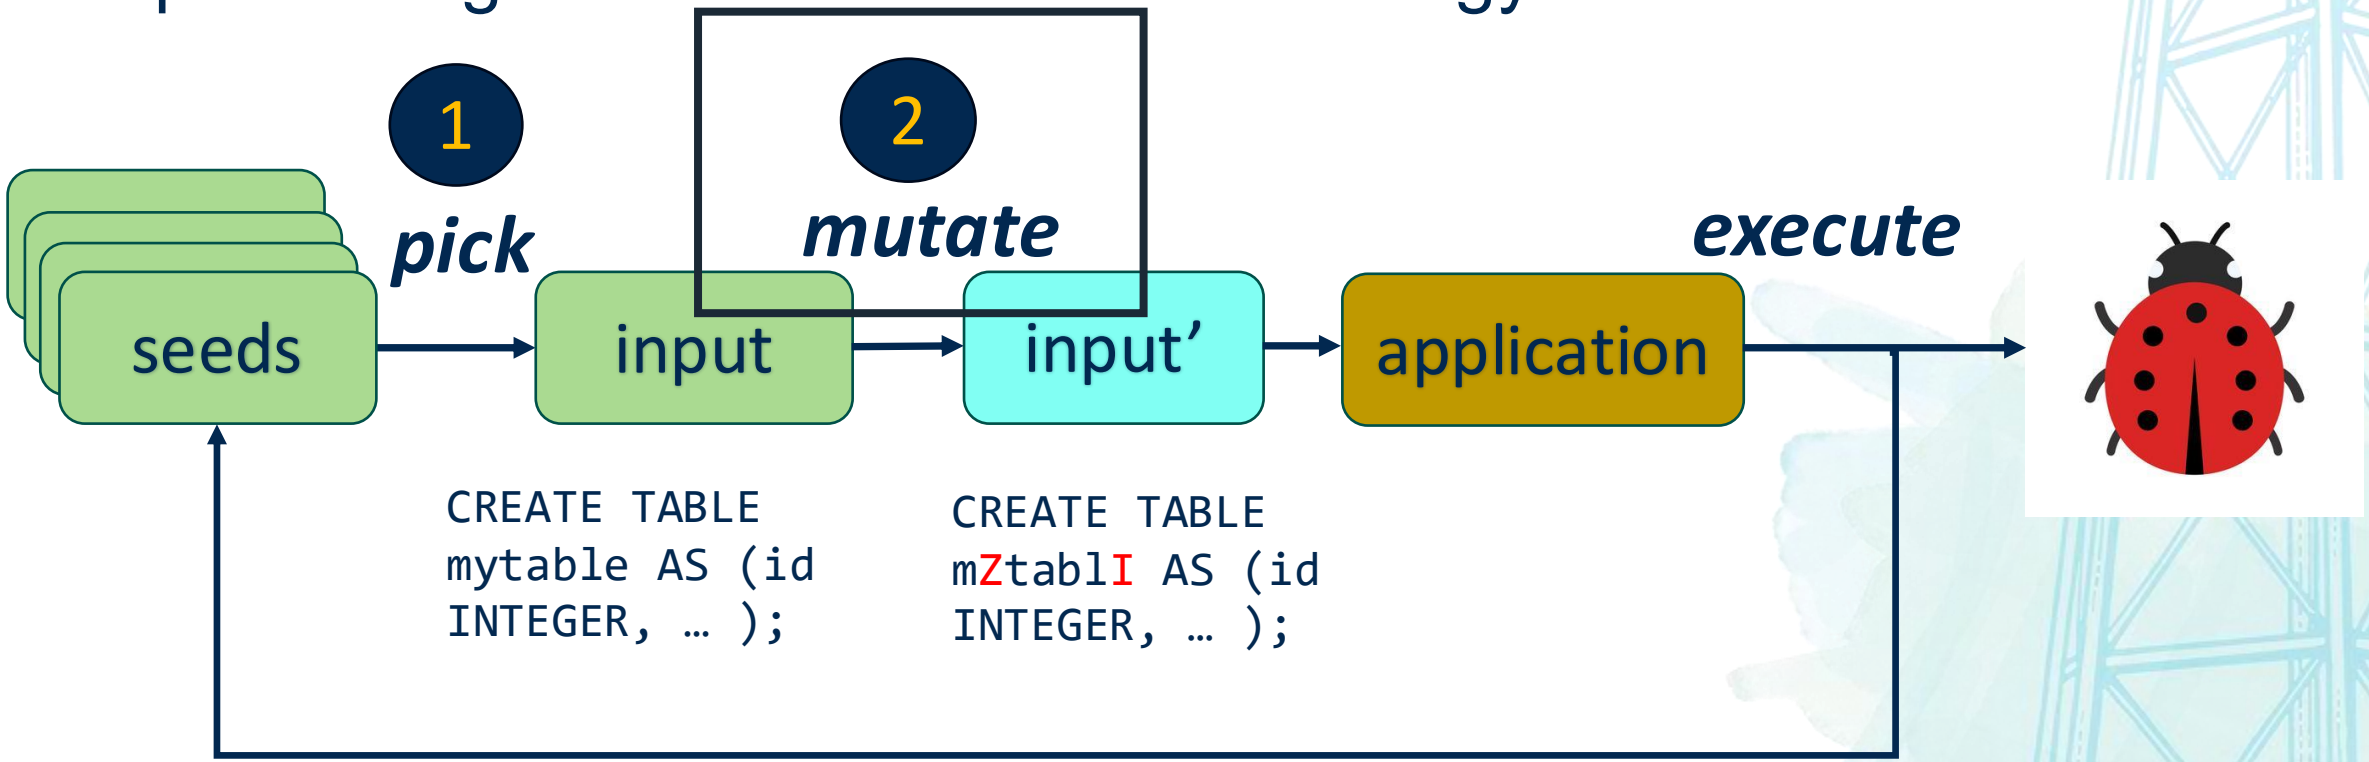

Mutation strategy – byte-level transformations

- Shuffle bytes
- Erase bytes
- Insert byte
- Change byte

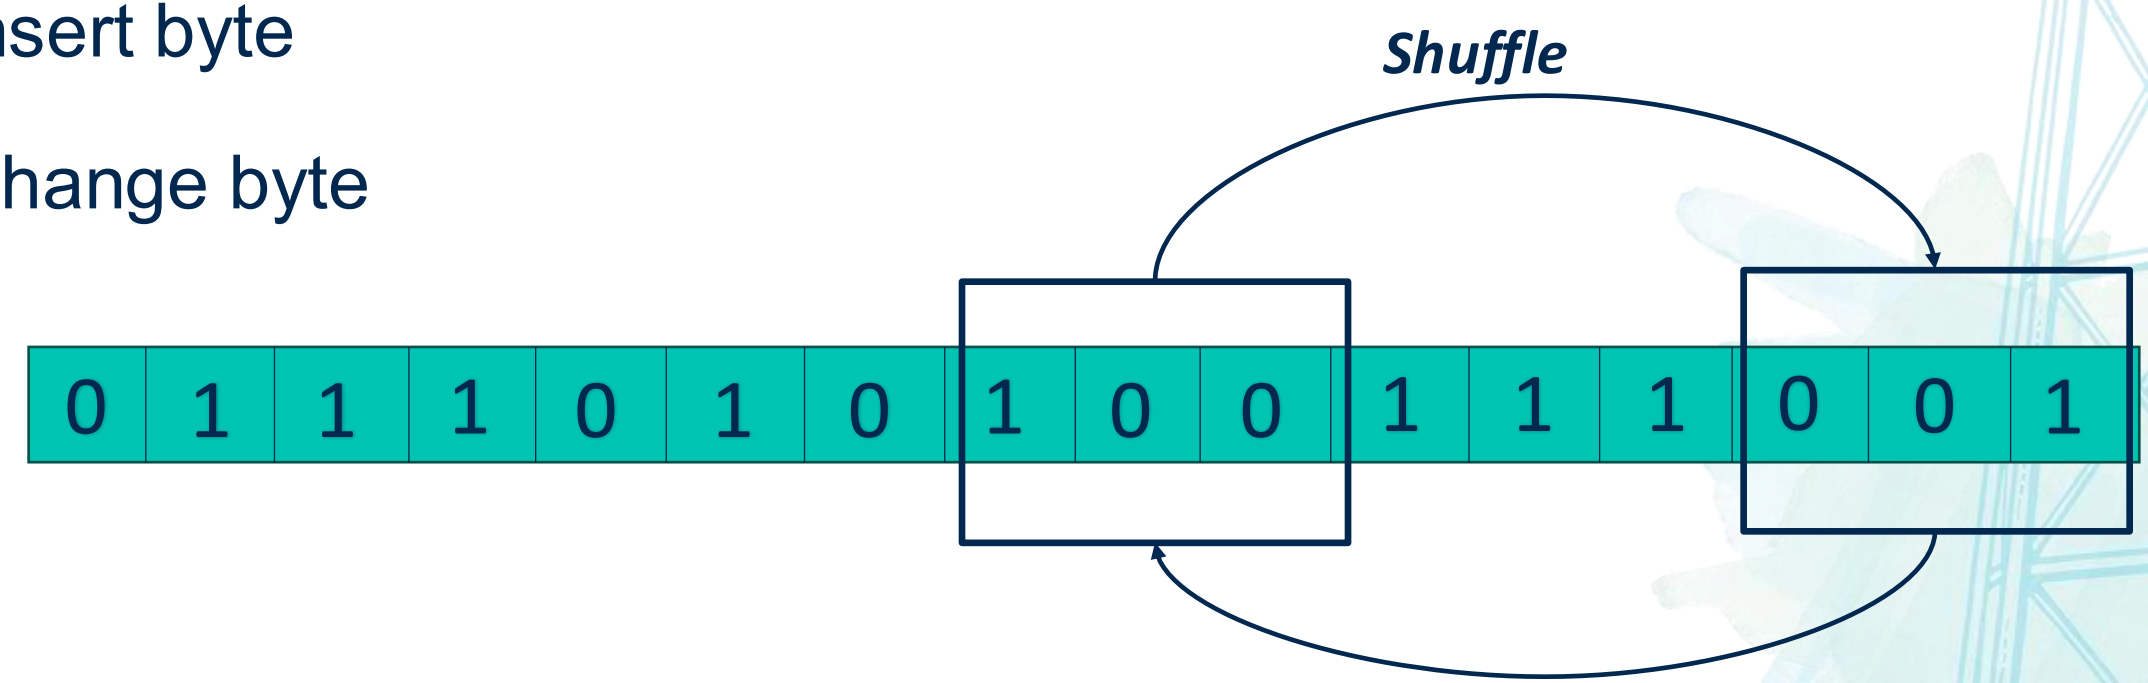

Mutation strategy – byte-level transformations

- Shuffle bytes
- Erase bytes
- Insert byte
- Change byte

New input

0	1	1	1	0	1	0	0	0	1	1	1	1	1	0	0

Mutation strategy – byte-level transformations

- Shuffle bytes
- Erase bytes
- Insert byte
- Change byte

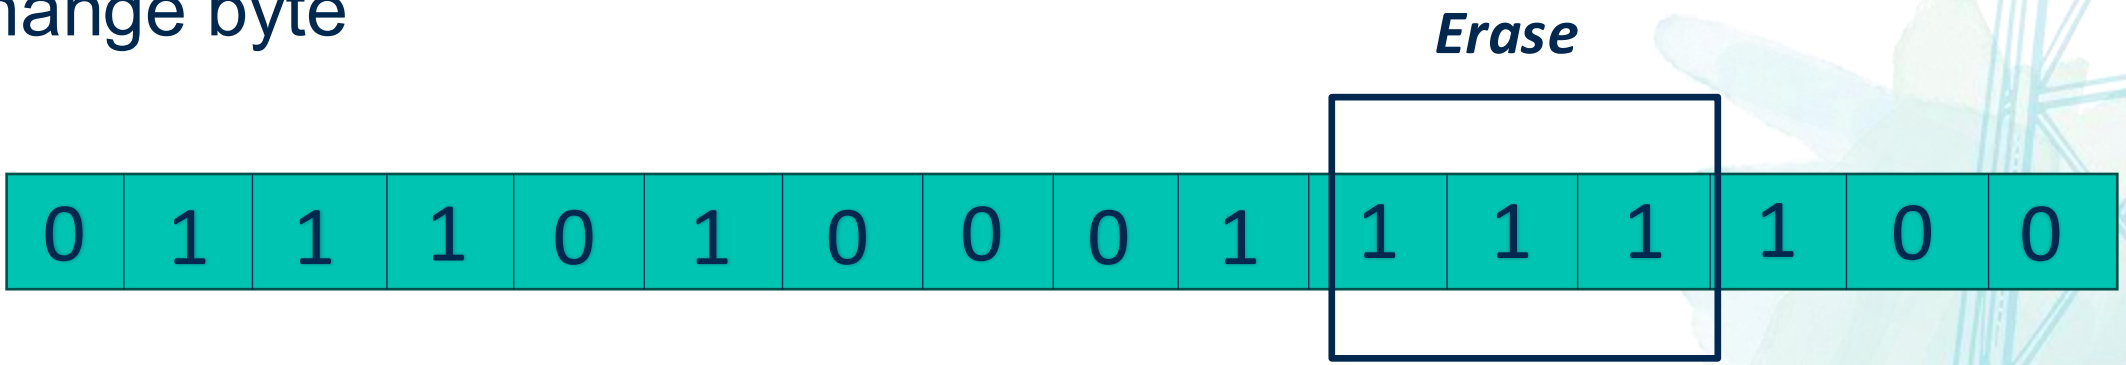

Mutation strategy – byte-level transformations

- Shuffle bytes
- Erase bytes
- Insert byte
- Change byte

Next input

0	1	1	1	0	1	0	0	0	1	1	0	0

Mutation strategy – byte-level transformations

- Benefits:
  - Simple to implement
  - Good for binary inputs and streams
    - JPEG file encoder/decoder
- Disadvantages:
  - Generates many invalid inputs for applications that expect structured inputs

Mutation strategy – dictionaries

- Programmer provides a dictionary of “keywords”
  - For SQL, could include `CREATE`, `DELETE`, `INSERT`, `INTO`, `FROM`, `INTEGER`, and so on…
  - For HTML, could include HTML tags `<html>`, `<body>`, `<p>`, and so on…
- Mutate the input by replacing or adding a chosen keyword from this dictionary
- Input - `<html><body><p>Hello world </p></body></html>`
  Dictionary - `{<html>, <body>, <p>,…}`
  Mutated input –
  `<html><body><p> Hello world <p>, Welcome!</p> </p></html>`

Mutation strategy – dictionaries

- Benefits:
  - Still simple to implement
  - Works reasonably well for structured inputs
- Disadvantages
  - Still generates many invalid inputs for applications that expect structured inputs

Characteristics

- All mutators produce ***many*** invalid inputs
- E.g. malformed HTML, malformed SQL queries
  - `<html><body></P.</ht~l>`
- Still produces enough valid test-cases to be ***very*** effective

Fuzzing pipeline

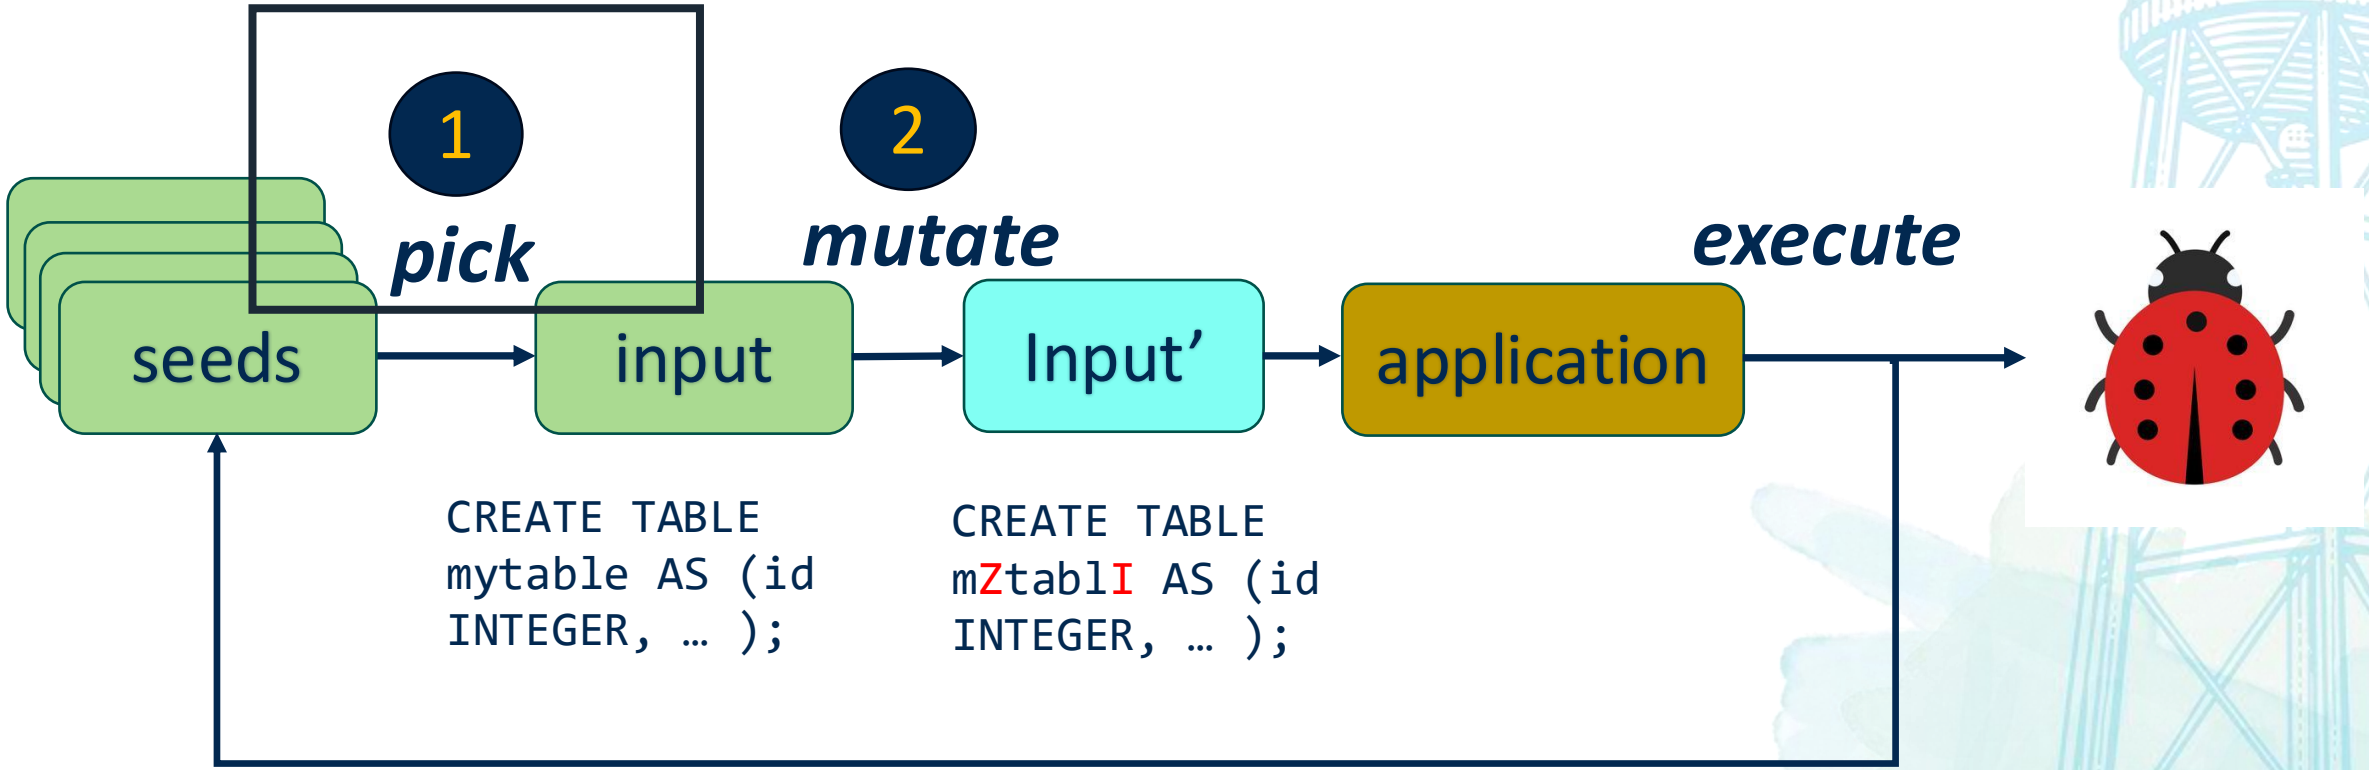

How to pick which seeds to mutate?

- Goal: Should pick seeds that are “most promising” and likely to find bugs
- What is ***“most promising”?***

Metric of seed “goodness”: coverage

- Amount of code or execution paths explored during fuzz testing
- Input seeds that increase coverage are “more promising”
  - Explore new program behaviors
  - Ensures untested paths are tested
- High coverage = high bug-finding potential

Metric of seed “goodness”: coverage

- Amount of code or execution paths explored during fuzz testing

```
int function(int a, int b) {
    if (a > 0) {
        if (b > 0) {
            if (b == 10) {
                printf(“Hello!\n”);
            } else {
                printf(“Goodbye!\n”);
            }
        } else {
            printf(“Welcome!\n”);
        }
    } else {
        printf(“Hi!\n”);
    }
}
```

Metric of seed “goodness”: coverage

- Amount of code or execution paths explored during fuzz testing
- Input: a = 10, b = 0 -> new coverage!

```
int function(int a, int b) {
    if (a > 0) {
        if (b > 0) {
            if (b == 10) {
                printf(“Hello!\n”);
            } else {
                printf(“Goodbye!\n”);
            }
        } else {
            printf(“Welcome!\n”);
        }
    } else {
        printf(“Hi!\n”);
    }
}
```

Metric of seed “goodness”: coverage

- Amount of code or execution paths explored during fuzz testing
- Input: a = 11, b = 0 -> no new coverage

```
int function(int a, int b) {
    if (a > 0) {
        if (b > 0) {
            if (b == 10) {
                printf(“Hello!\n”);
            } else {
                printf(“Goodbye!\n”);
            }
        } else {
            printf(“Welcome!\n”);
        }
    } else {
        printf(“Hi!\n”);
    }
}
```

Metric of seed “goodness”: coverage

- Amount of code or execution paths explored during fuzz testing
- Input: a = 10, b = 1 -> new coverage!

```
int function(int a, int b) {
    if (a > 0) {
        if (b > 0) {
            if (b == 10) {
                printf(“Hello!\n”);
            } else {
                printf(“Goodbye!\n”);
            }
        } else {
            printf(“Welcome!\n”);
        }
    } else {
        printf(“Hi!\n”);
    }
}
```

Metric of seed “goodness”: coverage

- Amount of code or execution paths explored during fuzz testing
- Input: a = 10, b = 2 -> no new coverage!

```
int function(int a, int b) {
    if (a > 0) {
        if (b > 0) {
            if (b == 10) {
                printf(“Hello!\n”);
            } else {
                printf(“Goodbye!\n”);
            }
        } else {
            printf(“Welcome!\n”);
        }
    } else {
        printf(“Hi!\n”);
    }
}
```

Metric of seed “goodness”: coverage

- Amount of code or execution paths explored during fuzz testing
- Input: a = 10, b = 3 -> no new coverage!

```
int function(int a, int b) {
    if (a > 0) {
        if (b > 0) {
            if (b == 10) {
                printf(“Hello!\n”);
            } else {
                printf(“Goodbye!\n”);
            }
        } else {
            printf(“Welcome!\n”);
        }
    } else {
        printf(“Hi!\n”);
    }
}
```

Metric of seed “goodness”: coverage

- Amount of code or execution paths explored during fuzz testing
- Input: a = -1, b = 3 -> new coverage!

```
int function(int a, int b) {
    if (a > 0) {
        if (b > 0) {
            if (b == 10) {
                printf(“Hello!\n”);
            } else {
                printf(“Goodbye!\n”);
            }
        } else {
            printf(“Welcome!\n”);
        }
    } else {
        printf(“Hi!\n”);
    }
}
```

Coverage guided fuzzing

- Record the coverage
- Prioritize mutating input seeds which give higher coverage

Coverage-guided fuzzer pipeline

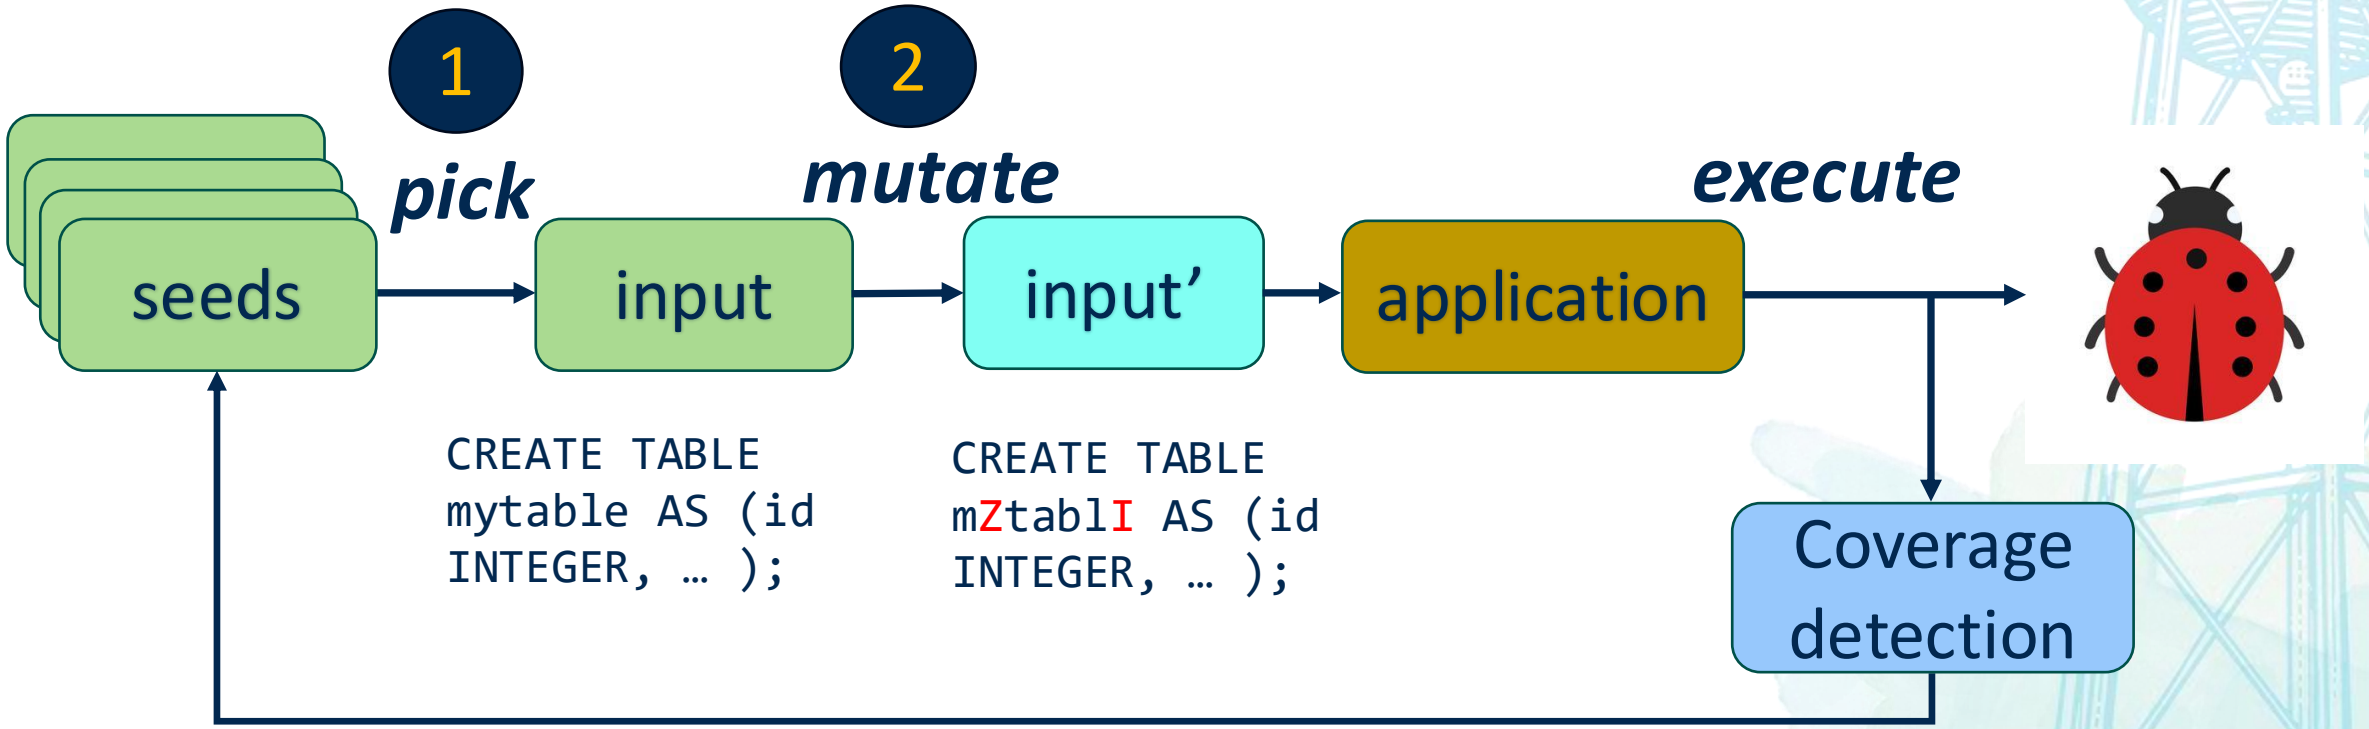

Add* input' *to* seeds *if new coverage observed

Fuzzing – high level ideas

- Randomly generate inputs
- Mutation fuzzing
  - Start with known valid inputs
  - Mutate them
    - Byte-level transformations
    - Dictionary-based transformations
- Pick inputs which resulted in new branch coverage for further mutation

AFL demo

- American fuzzy lop (AFL)
- Fuzzer for C/C++ applications

AFL demo

- https://github.com/wolframroesler/afl-demo

AFL demo

- Throughput – number of inputs tested per second
- ***Is higher always better?***

```
                    american fuzzy lop 2.52b (afldemo)

┌─ process timing ─────────────────────────────────────┬─ overall results ─────┐
│        run time : 0 days, 0 hrs, 0 min, 46 sec       │  cycles done : 0      │
│   last new path : 0 days, 0 hrs, 0 min, 0 sec        │  total paths : 309    │
│ last uniq crash : 0 days, 0 hrs, 0 min, 0 sec        │ uniq crashes : 215    │
│  last uniq hang : none seen yet                      │   uniq hangs : 0      │
├─ cycle progress ────────────────────┬─ map coverage ─┴───────────────────────┤
│  now processing : 27 (8.74%)        │    map density : 0.06% / 2.23%         │
│ paths timed out : 0 (0.00%)         │ count coverage : 1.08 bits/tuple       │
├─ stage progress ────────────────────┼─ findings in depth ────────────────────┤
│  now trying : havoc                 │ favored paths : 107 (34.63%)           │
│ stage execs : 12.9k/32.8k (39.37%)  │  new edges on : 287 (92.88%)           │
│ total execs : 128k                  │ total crashes : 2882 (215 unique)      │
│  exec speed : 2142/sec              │  total tmouts : 0 (0 unique)           │
├─ fuzzing strategy yields ───────────┴───────────────┬─ path geometry ────────┤
│   bit flips : 24/2024, 22/2016, 20/2000             │    levels : 4          │
│  byte flips : 3/253, 1/245, 0/231                   │   pending : 302        │
│ arithmetics : 167/14.2k, 0/1532, 0/140              │  pend fav : 101        │
│  known ints : 17/1368, 6/6782, 2/10.1k              │ own finds : 306        │
│  dictionary : 0/0, 0/0, 0/0                         │  imported : n/a        │
│       havoc : 193/71.9k, 0/0                        │ stability : 100.00%    │
│        trim : 0.79%/86, 0.00%                       ├────────────────────────┘
└─────────────────────────────────────────────────────┘          [cpu000: 31%]
```

AFL demo

- Crashes found by the fuzzer

```
                    american fuzzy lop 2.52b (afldemo)

┌─ process timing ─────────────────────────────────────┬─ overall results ─────┐
│        run time : 0 days, 0 hrs, 0 min, 46 sec       │  cycles done : 0      │
│   last new path : 0 days, 0 hrs, 0 min, 0 sec        │  total paths : 309    │
│ last uniq crash : 0 days, 0 hrs, 0 min, 0 sec        │ uniq crashes : 215    │
│  last uniq hang : none seen yet                      │   uniq hangs : 0      │
├─ cycle progress ────────────────────┬─ map coverage ─┴───────────────────────┤
│  now processing : 27 (8.74%)        │    map density : 0.06% / 2.23%         │
│ paths timed out : 0 (0.00%)         │ count coverage : 1.08 bits/tuple       │
├─ stage progress ────────────────────┼─ findings in depth ────────────────────┤
│  now trying : havoc                 │ favored paths : 107 (34.63%)           │
│ stage execs : 12.9k/32.8k (39.37%)  │  new edges on : 287 (92.88%)           │
│ total execs : 128k                  │ total crashes : 2882 (215 unique)      │
│  exec speed : 2142/sec              │  total tmouts : 0 (0 unique)           │
├─ fuzzing strategy yields ───────────┴───────────────┬─ path geometry ────────┤
│   bit flips : 24/2024, 22/2016, 20/2000             │    levels : 4          │
│  byte flips : 3/253, 1/245, 0/231                   │   pending : 302        │
│ arithmetics : 167/14.2k, 0/1532, 0/140              │  pend fav : 101        │
│  known ints : 17/1368, 6/6782, 2/10.1k              │ own finds : 306        │
│  dictionary : 0/0, 0/0, 0/0                         │  imported : n/a        │
│       havoc : 193/71.9k, 0/0                        │ stability : 100.00%    │
│        trim : 0.79%/86, 0.00%                       ├────────────────────────┘
└─────────────────────────────────────────────────────┘          [cpu000: 31%]
```

Crashes

AFL demo

- Coverage – number of code paths executed

```
                    american fuzzy lop 2.52b (afldemo)

┌─ process timing ─────────────────────────────────────┬─ overall results ─────┐
│        run time : 0 days, 0 hrs, 0 min, 46 sec       │  cycles done : 0      │
│   last new path : 0 days, 0 hrs, 0 min, 0 sec        │  total paths : 309    │
│ last uniq crash : 0 days, 0 hrs, 0 min, 0 sec        │ uniq crashes : 215    │
│  last uniq hang : none seen yet                      │   uniq hangs : 0      │
├─ cycle progress ────────────────────┬─ map coverage ─┴───────────────────────┤
│  now processing : 27 (8.74%)        │    map density : 0.06% / 2.23%         │
│ paths timed out : 0 (0.00%)         │ count coverage : 1.08 bits/tuple       │
├─ stage progress ────────────────────┼─ findings in depth ────────────────────┤
│  now trying : havoc                 │ favored paths : 107 (34.63%)           │
│ stage execs : 12.9k/32.8k (39.37%)  │  new edges on : 287 (92.88%)           │
│ total execs : 128k                  │ total crashes : 2882 (215 unique)      │
│  exec speed : 2142/sec              │  total tmouts : 0 (0 unique)           │
├─ fuzzing strategy yields ───────────┴───────────────┬─ path geometry ────────┤
│   bit flips : 24/2024, 22/2016, 20/2000             │    levels : 4          │
│  byte flips : 3/253, 1/245, 0/231                   │   pending : 302        │
│ arithmetics : 167/14.2k, 0/1532, 0/140              │  pend fav : 101        │
│  known ints : 17/1368, 6/6782, 2/10.1k              │ own finds : 306        │
│  dictionary : 0/0, 0/0, 0/0                         │  imported : n/a        │
│       havoc : 193/71.9k, 0/0                        │ stability : 100.00%    │
│        trim : 0.79%/86, 0.00%                       ├────────────────────────┘
└─────────────────────────────────────────────────────┘          [cpu000: 31%]
```

Coverage

Fuzz testing - sanitizers

Tapti Palit

Recall: memory safety bugs

- Fuzzers execute the program with random inputs and look for crashes
- But memory safety bugs may or may not crash the program
- **Undefined Behavior (UB)**

Sanitizers: oracles for memory safety bugs

- Adds additional compiler instrumentation to record object metadata
- Turns silent UBs to crashes
- Popular sanitizers
  - Address sanitizer (ASAN) – can detect buffer overflow
  - Leak sanitizer (LSAN) – can detect memory leak
  - Thread sanitizer (TSAN) – can detect data races
- ***Significantly*** slows down the application; but can detect memory safety bugs!!

Address sanitizer

- Goal – detect buffer overflow and UAF
- Insert “red/poison zones” around memory
- Updates to red zones indicate buffer overflow
  - Need to detect writes to red zones

```
void function(char* p) {
    while(*p != ‘\0’) *p++ = ‘a’;
}


int main(void) {
    char* s = malloc(16); // heap object 1
    char* p = malloc(8); // heap object 2
    function(s);
}
```

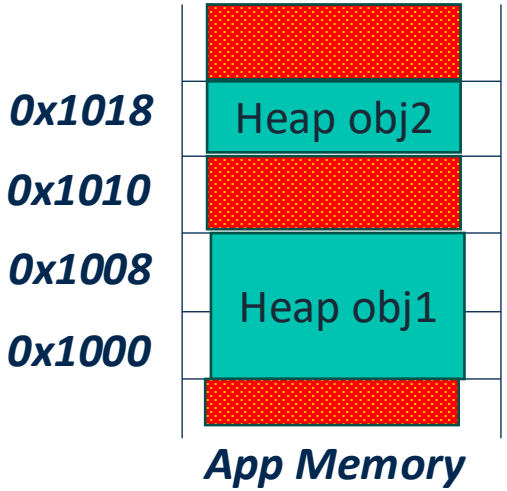

Shadow memory

- Maintains shadow memory to record metadata about red and accessible zones
- Each ***byte*** in shadow memory represents a corresponding ***8 bytes*** in app
  - 0x00 for valid
  - 0xFF for poisoned range
  - 0x01 – 0x07 – partially poisoned

```
void function(char* p) {
     while(*p != '\0') *p++ = 'a';
}


int main(void) {
    char* s = malloc(16); // heap object 1
    char* p = malloc(8); // heap object 2
    function(s);
}
```

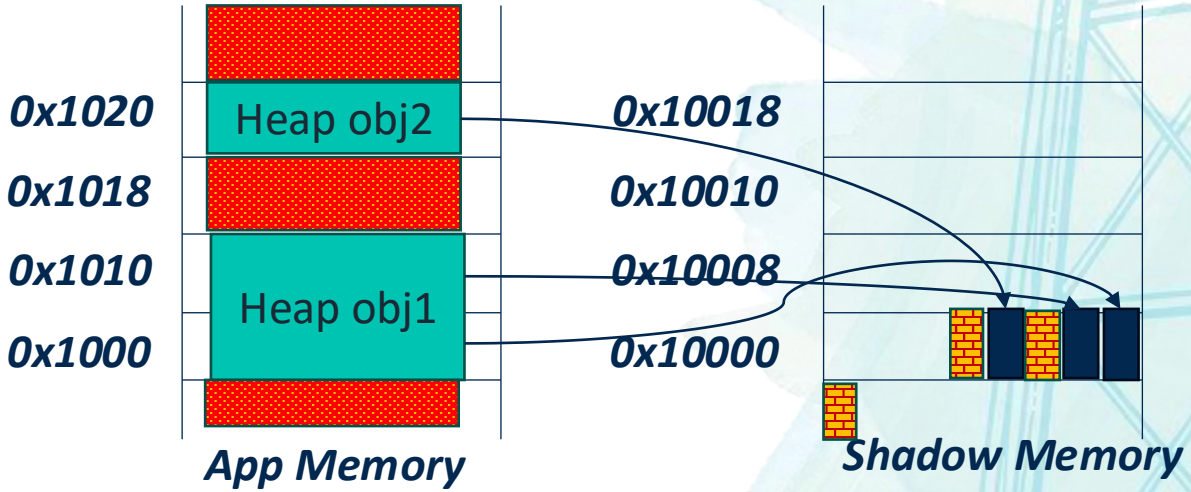

Compiler checks

- Compiler adds instructions before each memory access to first check the shadow memory byte

```
void function(char* p) {
     while(*p != '\0') *p++ = 'a';
}


int main(void) {
    char* s = malloc(16); // heap object 1
    char* p = malloc(8); // heap object 2
    function(s);
}
```

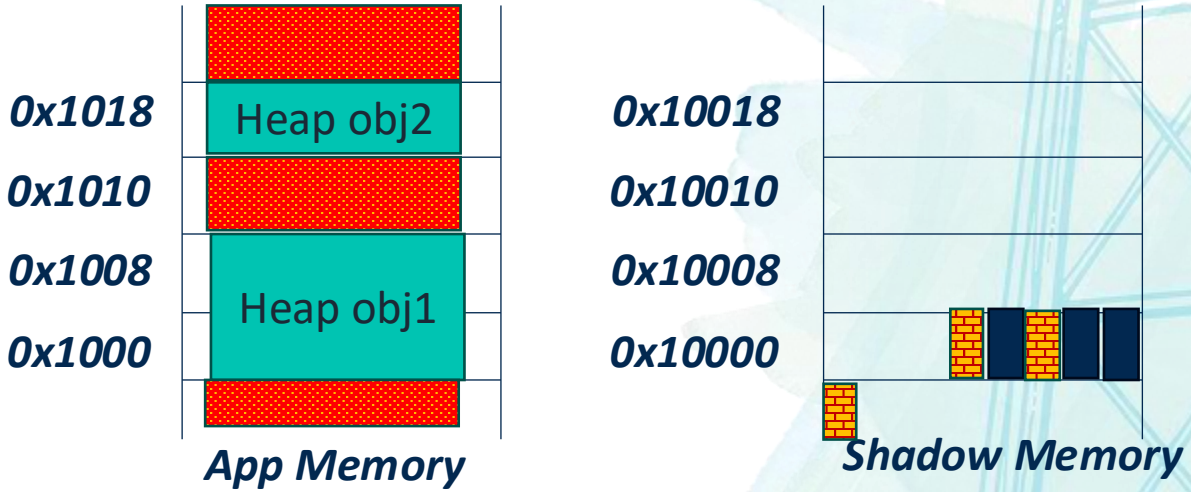

Compiler checks pseudocode

- Compiler adds instructions before each memory access to first check the shadow memory byte
- Pseudocode
  - shdw_poisoned(void*) computes the address of the shadow byte
  - Checks if it is fully/partially poisoned

```
void function(char* p) {
     while(*p != '\0') {
        if (shdw_poisoned(p)) { assert(false); }
        *p++ = 'a';
     }
}

int main(void) {
    char* s = malloc(16); // heap object 1
    char* p = malloc(8); // heap object 2
    function(s);
}
```

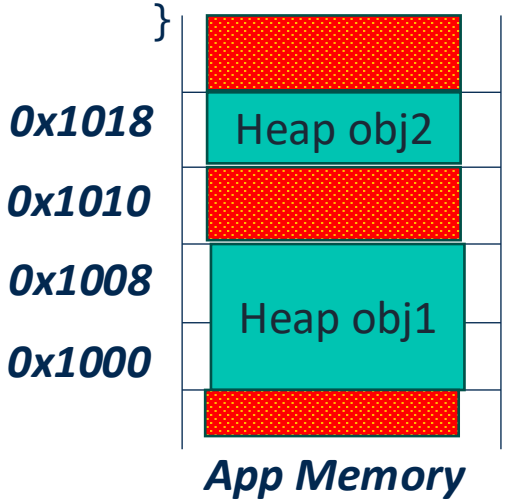

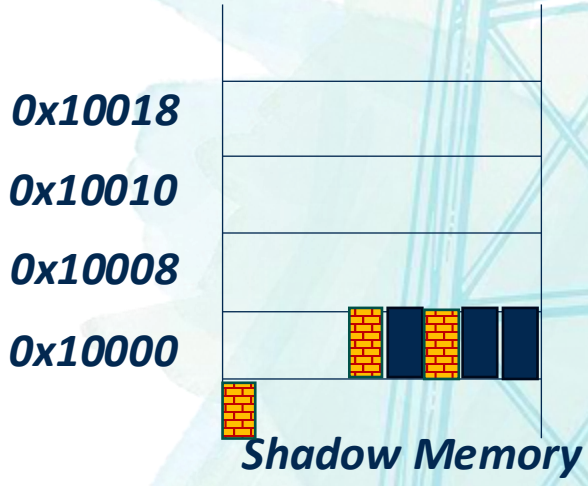

Why shadow memory?

- Shadow memory stores metadata about the 8-byte memory range
- Fast computation of shadow byte address
  - `shadow = (address >> 3) + SHADOW_OFFSET`
- Fast checks of shadow byte value
  - `shadow_value = *shadow`

Summary

- Fuzzing is an effective testing mechanism
- Unlike other testing approaches it can ***find new inputs*** that cause program misbehavior
- Mutation and coverage are important for generating effective fuzzers
- Sanitizers provide additional information needed to find *silent UB* bugs